

# Agentic Workflows

## A Practical Guide

# Agentic Workflows

A Practical Guide

Agentbook Contributors

February 2026

# Contents

|   |           |
|---|-----------|
| <b>Copyright</b>                                | <b>10</b> |
| <b>Preface</b>                                  | <b>11</b> |
| <b>Conventions Used in This Book</b>            | <b>12</b> |
| <b>1 Introduction to Agentic Workflows</b>      | <b>13</b> |
| 1.1 Chapter Preview . . . . .                   | 13        |
| 1.2 What Are Agentic Workflows?                 | 13        |
| 1.2.1 Key Concepts . . . . .                    | 13        |
| 1.2.2 Terminology and Roles . . . . .           | 13        |
| 1.2.3 Why Agentic Workflows? . . . . .          | 14        |
| 1.3 Real-World Applications . . . . .           | 14        |
| 1.3.1 Software Development . . . . .            | 14        |
| 1.3.2 Content Management . . . . .              | 14        |
| 1.3.3 Operations . . . . .                      | 14        |
| 1.4 The Agent Development Lifecycle . . . . .   | 14        |
| 1.5 Getting Started . . . . .                   | 14        |
| 1.6 Key Takeaways . . . . .                     | 14        |
| <b>2 Agent Orchestration</b>                    | <b>15</b> |
| 2.1 Chapter Preview . . . . .                   | 15        |
| 2.2 Understanding Agent Orchestration . . . . . | 15        |
| 2.3 Orchestration Patterns . . . . .            | 15        |
| 2.3.1 Sequential Execution . . . . .            | 15        |
| 2.3.2 Parallel Execution . . . . .              | 15        |
| 2.3.3 Hierarchical Execution . . . . .          | 15        |
| 2.3.4 Event-Driven Orchestration . . . . .      | 16        |
| 2.4 Coordination Mechanisms . . . . .           | 16        |
| 2.4.1 Message Passing . . . . .                 | 16        |
| 2.4.2 Shared State . . . . .                    | 16        |
| 2.4.3 Direct Invocation . . . . .               | 16        |
| 2.5 Best Practices . . . . .                    | 16        |
| 2.5.1 Clear Responsibilities . . . . .          | 16        |
| 2.5.2 Error Handling . . . . .                  | 16        |
| 2.5.3 Monitoring . . . . .                      | 16        |
| 2.5.4 Isolation . . . . .                       | 17        |
| 2.6 Key Takeaways . . . . .                     | 17        |
| 2.7 Orchestration Frameworks . . . . .          | 17        |
| 2.7.1 GitHub Actions . . . . .                  | 17        |
| 2.7.2 LangChain . . . . .                       | 17        |
| 2.7.3 Custom Orchestration . . . . .            | 17        |

|          |  |           |
|----------|--|-----------|
| 2.8      | Real-World Example: Self-Updating Documentation . . . . .            | 18        |
| 2.9      | Challenges and Solutions . . . . .                                   | 18        |
| 2.9.1    | Challenge: Agent Conflicts . . . . .                                 | 18        |
| 2.9.2    | Challenge: Debugging . . . . .                                       | 18        |
| 2.9.3    | Challenge: Performance . . . . .                                     | 18        |
| 2.9.4    | Challenge: Versioning . . . . .                                      | 18        |
| 2.10     | Key Takeaways . . . . .  | 18        |
| <b>3</b> | <b>Agentic Scaffolding</b>   | <b>19</b> |
| 3.1      | Chapter Preview . . . . .  | 19        |
| 3.2      | What Is Agentic Scaffolding? . . . . .                               | 19        |
| 3.3      | Core Components . . . . .  | 19        |
| 3.3.1    | Tool Access Layer . . . . .  | 19        |
| 3.3.2    | Context Management . . . . .   | 20        |
| 3.3.3    | Execution Environment . . . . .                                      | 20        |
| 3.3.4    | Communication Protocol . . . . .                                     | 20        |
| 3.4      | Scaffolding Patterns . . . . .                                       | 21        |
| 3.4.1    | Pattern 1: Tool Composition . . . . .                                | 21        |
| 3.4.2    | Pattern 2: Skill Libraries . . . . .                                 | 21        |
| 3.4.3    | Pattern 3: Resource Management . . . . .                             | 22        |
| 3.4.4    | Pattern 4: Observability . . . . .                                   | 22        |
| 3.5      | Building Scaffolding: Step by Step . . . . .                         | 23        |
| 3.5.1    | Step 1: Define Your Agent Ecosystem . . . . .                        | 23        |
| 3.5.2    | Step 2: Implement Tool Registry . . . . .                            | 23        |
| 3.5.3    | Step 3: Create Agent Templates . . . . .                             | 23        |
| 3.5.4    | Step 4: Implement Error Recovery . . . . .                           | 24        |
| 3.6      | Scaffolding for This Book . . . . .                                  | 24        |
| 3.7      | Best Practices . . . . .   | 24        |
| 3.8      | Common Pitfalls . . . . .  | 24        |
| 3.9      | Key Takeaways . . . . .  | 24        |
| <b>4</b> | <b>Skills and Tools Management</b>                                   | <b>26</b> |
| 4.1      | Chapter Preview . . . . .  | 26        |
| 4.2      | Understanding Skills vs. Tools . . . . .                             | 26        |
| 4.2.1    | Tools . . . . .  | 26        |
| 4.2.2    | Skills . . . . .   | 26        |
| 4.3      | Tool Design Principles . . . . .                                     | 26        |
| 4.3.1    | Single Responsibility . . . . .                                      | 26        |
| 4.3.2    | Clear Interfaces . . . . .   | 27        |
| 4.3.3    | Error Handling . . . . .   | 27        |
| 4.3.4    | Documentation . . . . .  | 28        |
| 4.4      | Creating Custom Tools . . . . .                                      | 28        |
| 4.4.1    | Basic Tool Template . . . . .  | 28        |
| 4.4.2    | Example: Markdown Validation Tool . . . . .                          | 29        |
| 4.5      | Agent Skills Standard (Primary Reference) . . . . .                  | 30        |
| 4.5.1    | Canonical Layout . . . . .   | 30        |
| 4.5.2    | Conformance Notes: Agent Skills vs. JSON-RPC Runtime Specs . . . . . | 31        |
| 4.5.3    | Relationship to MCP . . . . .  | 31        |
| 4.6      | Skill Development . . . . .  | 32        |
| 4.6.1    | Skill Architecture . . . . .   | 32        |
| 4.6.2    | Example: Code Review Skill . . . . .                                 | 32        |
| 4.7      | Importing and Using Skills . . . . .                                 | 33        |
| 4.7.1    | Skill Registry . . . . .   | 33        |
| 4.7.2    | Skill Composition . . . . .  | 34        |

|          |   |           |
|----------|---|-----------|
| 4.8      | Tool Discovery and Documentation . . . . .                      | 35        |
| 4.8.1    | Self-Documenting Tools . . . . .                                | 35        |
| 4.9      | Integrations: Connecting Tools to Real-World Surfaces . . . . . | 36        |
| 4.10     | Case Study: OpenClaw and pi-mono . . . . .                      | 36        |
| 4.10.1   | OpenClaw Architecture in Detail . . . . .                       | 36        |
| 4.10.2   | Key Design Principles . . . . .                                 | 37        |
| 4.11     | Related Architectures and Frameworks . . . . .                  | 37        |
| 4.11.1   | LangChain and LangGraph . . . . .                               | 37        |
| 4.11.2   | CrewAI . . . . .  | 38        |
| 4.11.3   | Microsoft Semantic Kernel . . . . .                             | 38        |
| 4.11.4   | AutoGen . . . . .   | 38        |
| 4.11.5   | Comparing Architecture Patterns . . . . .                       | 39        |
| 4.12     | MCP: Modern Tooling and Adoption . . . . .                      | 39        |
| 4.12.1   | What MCP Brings to Tools . . . . .                              | 39        |
| 4.12.2   | Common Usage Patterns . . . . .                                 | 39        |
| 4.12.3   | Acceptance Across Major Clients . . . . .                       | 40        |
| 4.12.4   | Practical Guidance for Authors and Teams . . . . .              | 40        |
| 4.13     | Best Practices . . . . .  | 40        |
| 4.13.1   | Version Tools and Skills . . . . .                              | 40        |
| 4.13.2   | Test Independently . . . . .                                    | 40        |
| 4.13.3   | Provide Fallbacks . . . . .                                     | 41        |
| 4.13.4   | Monitor Usage . . . . .   | 41        |
| 4.14     | Emerging Standards: AGENTS.md . . . . .                         | 41        |
| 4.14.1   | The AGENTS.md Pseudo-Standard . . . . .                         | 41        |
| 4.14.2   | Related Standards: Skills and Capabilities . . . . .            | 42        |
| 4.15     | How Agents Become Aware of Imports . . . . .                    | 43        |
| 4.15.1   | The Import Awareness Problem . . . . .                          | 43        |
| 4.15.2   | Mechanisms for Import Discovery . . . . .                       | 43        |
| 4.15.3   | Best Practices for Import-Aware Agents . . . . .                | 46        |
| 4.15.4   | Alias Conventions . . . . .                                     | 46        |
| 4.15.5   | Import Awareness in Multi-Agent Systems . . . . .               | 48        |
| 4.16     | Key Takeaways . . . . .   | 48        |
| <b>5</b> | <b>Discovery and Imports</b>                                    | <b>50</b> |
| 5.1      | Chapter Preview . . . . .                                       | 50        |
| 5.2      | Terminology Baseline for the Rest of the Book . . . . .         | 50        |
| 5.3      | A Taxonomy That Disambiguates What We Are Discovering . . . . . | 50        |
| 5.3.1    | 1) Tool artefacts . . . . .                                     | 50        |
| 5.3.2    | 2) Skill artefacts . . . . .                                    | 50        |
| 5.3.3    | 3) Agent artefacts . . . . .                                    | 51        |
| 5.3.4    | 4) Workflow-fragment artefacts . . . . .                        | 51        |
| 5.3.5    | Confusing cases to stop using . . . . .                         | 51        |
| 5.4      | Discovery Mechanisms . . . . .                                  | 51        |
| 5.4.1    | Local scan . . . . .  | 51        |
| 5.4.2    | Registry discovery . . . . .                                    | 51        |
| 5.4.3    | Domain-convention discovery . . . . .                           | 51        |
| 5.4.4    | Explicit configuration . . . . .                                | 51        |
| 5.5      | Import, Install, Activate: Three Different Operations . . . . . | 52        |
| 5.5.1    | Import . . . . .  | 52        |
| 5.5.2    | Install . . . . .   | 52        |
| 5.5.3    | Activate . . . . .  | 52        |
| 5.6      | Trust Boundaries and Supply Chain (Compact Model) . . . . .     | 52        |
| 5.7      | Worked Examples . . . . .                                       | 52        |
| 5.7.1    | Example A: GH-AW workflow-fragment import . . . . .             | 52        |

|          |  |           |
|----------|--|-----------|
| 5.7.2    | Example B: AGENTS.md import conventions as resolution policy . . . . . | 53        |
| 5.7.3    | Example C: Capability discovery and activation policy . . . . .        | 53        |
| 5.8      | Key Takeaways . . . . .  | 53        |
| <b>6</b> | <b>GitHub Agentic Workflows (GH-AW)</b>                                | <b>54</b> |
| 6.1      | Chapter Preview . . . . .  | 54        |
| 6.2      | Why GH-AW Matters . . . . .  | 54        |
| 6.3      | Core Workflow Structure . . . . .                                      | 54        |
| 6.4      | How GH-AW Runs . . . . .   | 55        |
| 6.4.1    | File Location . . . . .  | 55        |
| 6.4.2    | Key Behaviors . . . . .  | 55        |
| 6.5      | Compilation Model Examples . . . . .                                   | 55        |
| 6.5.1    | Example 1: Issue Triage Workflow . . . . .                             | 56        |
| 6.5.2    | Example 2: Reusable Component + Import . . . . .                       | 57        |
| 6.5.3    | Example 3: Safe Outputs in the Compiled Job . . . . .                  | 58        |
| 6.6      | Tools, Safe Inputs, and Safe Outputs . . . . .                         | 59        |
| 6.6.1    | Tools . . . . .  | 59        |
| 6.6.2    | Safe Outputs . . . . .   | 59        |
| 6.6.3    | Safe Inputs . . . . .  | 59        |
| 6.7      | Imports and Reusable Components . . . . .                              | 59        |
| 6.8      | ResearchPlanAssign: A Pattern for Self-Maintaining Books . . . . .     | 59        |
| 6.9      | Applying GH-AW to This Repository . . . . .                            | 60        |
| 6.10     | Key Takeaways . . . . .  | 60        |
| <b>7</b> | <b>GitHub Agents</b>   | <b>61</b> |
| 7.1      | Chapter Preview . . . . .  | 61        |
| 7.2      | Understanding GitHub Agents . . . . .                                  | 61        |
| 7.3      | The GitHub Agent Ecosystem . . . . .                                   | 61        |
| 7.3.1    | GitHub Copilot . . . . .   | 61        |
| 7.3.2    | GitHub Copilot Coding Agent . . . . .                                  | 62        |
| 7.3.3    | GitHub Actions Agents . . . . .  | 62        |
| 7.4      | Agent Capabilities . . . . .   | 62        |
| 7.4.1    | Reading and Understanding . . . . .                                    | 62        |
| 7.4.2    | Writing and Creating . . . . .   | 63        |
| 7.4.3    | Reasoning and Deciding . . . . .                                       | 63        |
| 7.5      | Multi-Agent Orchestration . . . . .                                    | 63        |
| 7.5.1    | Why Multiple Agents? . . . . .   | 63        |
| 7.5.2    | Orchestration Patterns . . . . .                                       | 63        |
| 7.5.3    | Agent Handoff Protocol . . . . .                                       | 64        |
| 7.6      | Implementing GitHub Agents . . . . .                                   | 64        |
| 7.6.1    | Agent Definition Files . . . . .                                       | 64        |
| 7.6.2    | Agent Configuration . . . . .  | 65        |
| 7.6.3    | Error Handling . . . . .   | 65        |
| 7.7      | Best Practices . . . . .   | 66        |
| 7.7.1    | Clear Agent Personas . . . . .   | 66        |
| 7.7.2    | Structured Communication . . . . .                                     | 66        |
| 7.7.3    | Human Checkpoints . . . . .  | 66        |
| 7.7.4    | Audit Trail . . . . .  | 66        |
| 7.7.5    | Graceful Degradation . . . . .   | 66        |
| 7.8      | Security Considerations . . . . .                                      | 67        |
| 7.8.1    | Least Privilege . . . . .  | 67        |
| 7.8.2    | Input Validation . . . . .   | 67        |
| 7.8.3    | Output Sanitization . . . . .  | 67        |
| 7.8.4    | Protected Resources . . . . .  | 67        |

|          |   |           |
|----------|---|-----------|
| 7.9      | Real-World Example: This Book . . . . .                   | 67        |
| 7.9.1    | The Multi-Agent Workflow . . . . .                        | 67        |
| 7.9.2    | How It Works . . . . .                                    | 68        |
| 7.9.3    | Configuration . . . . .                                   | 68        |
| 7.10     | Multi-Agent Platform Compatibility . . . . .              | 68        |
| 7.10.1   | The Challenge of Agent Diversity . . . . .                | 68        |
| 7.10.2   | Repository Documentation as Agent Configuration . . . . . | 68        |
| 7.10.3   | The copilot-instructions.md File . . . . .                | 69        |
| 7.10.4   | Cross-Platform Strategy . . . . .                         | 69        |
| 7.10.5   | This Repository's Approach . . . . .                      | 70        |
| 7.10.6   | Best Practices for Agent-Friendly Repositories . . . . .  | 70        |
| 7.11     | Future of GitHub Agents . . . . .                         | 70        |
| 7.11.1   | Emerging Capabilities . . . . .                           | 70        |
| 7.11.2   | Integration Trends . . . . .                              | 70        |
| 7.12     | Key Takeaways . . . . .                                   | 70        |
| 7.13     | Learn More . . . . .                                      | 71        |
| 7.13.1   | Repository Documentation . . . . .                        | 71        |
| 7.13.2   | Agent Configuration Files . . . . .                       | 71        |
| 7.13.3   | Related Sections . . . . .                                | 71        |
| <b>8</b> | <b>Agents for Coding</b>                                  | <b>72</b> |
| 8.1      | Chapter Preview . . . . .                                 | 72        |
| 8.2      | Introduction . . . . .                                    | 72        |
| 8.3      | The Evolution of Coding Agents . . . . .                  | 72        |
| 8.3.1    | From Autocomplete to Autonomy . . . . .                   | 72        |
| 8.3.2    | Current Capabilities . . . . .                            | 72        |
| 8.4      | Specialized Architectures . . . . .                       | 73        |
| 8.4.1    | Single-Agent Architectures . . . . .                      | 73        |
| 8.4.2    | Multi-Agent Architectures . . . . .                       | 73        |
| 8.4.3    | Subagent and Swarms Mode . . . . .                        | 74        |
| 8.5      | Scaffolding for Coding Agents . . . . .                   | 75        |
| 8.5.1    | Project Initialization . . . . .                          | 75        |
| 8.5.2    | The AGENTS.md Standard . . . . .                          | 75        |
| 8.5.3    | Context Management . . . . .                              | 76        |
| 8.5.4    | Tool Registries . . . . .                                 | 77        |
| 8.6      | Leading Coding Agent Platforms . . . . .                  | 78        |
| 8.6.1    | GitHub Copilot and Coding Agent . . . . .                 | 78        |
| 8.6.2    | Claude Code . . . . .                                     | 79        |
| 8.6.3    | Cursor AI . . . . .                                       | 79        |
| 8.6.4    | CodeGPT and Agent Marketplaces . . . . .                  | 79        |
| 8.7      | Best Practices . . . . .                                  | 79        |
| 8.7.1    | Clear Task Boundaries . . . . .                           | 79        |
| 8.7.2    | Incremental Changes . . . . .                             | 79        |
| 8.7.3    | Test-Driven Development . . . . .                         | 80        |
| 8.7.4    | Human Review Integration . . . . .                        | 81        |
| 8.8      | Common Challenges . . . . .                               | 81        |
| 8.8.1    | Context Window Limitations . . . . .                      | 81        |
| 8.8.2    | Hallucination and Accuracy . . . . .                      | 82        |
| 8.8.3    | Security Concerns . . . . .                               | 82        |
| 8.9      | Key Takeaways . . . . .                                   | 82        |
| <b>9</b> | <b>Agents for Mathematics and Physics</b>                 | <b>84</b> |
| 9.1      | Chapter Preview . . . . .                                 | 84        |
| 9.2      | Introduction . . . . .                                    | 84        |

|           |   |           |
|-----------|---|-----------|
| 9.3       | The Landscape of Scientific Agents . . . . .              | 84        |
| 9.3.1     | Distinct Requirements . . . . .                           | 84        |
| 9.3.2     | Categories of Scientific Agents . . . . .                 | 84        |
| 9.4       | Theorem Proving Agents . . . . .                          | 85        |
| 9.4.1     | Formal Verification Background . . . . .                  | 85        |
| 9.4.2     | Ax-Prover Architecture . . . . .                          | 85        |
| 9.4.3     | Integration with Proof Assistants . . . . .               | 86        |
| 9.4.4     | Challenges in Theorem Proving . . . . .                   | 86        |
| 9.5       | Symbolic Computation Agents . . . . .                     | 86        |
| 9.5.1     | Computer Algebra Systems . . . . .                        | 86        |
| 9.5.2     | Combining Symbolic and Neural Approaches . . . . .        | 88        |
| 9.6       | Physics Simulation Agents . . . . .                       | 88        |
| 9.6.1     | Computational Physics Workflows . . . . .                 | 88        |
| 9.6.2     | Quantum Physics Specialization . . . . .                  | 89        |
| 9.7       | Scaffolding for Scientific Agents . . . . .               | 90        |
| 9.7.1     | Tool Integration Layer . . . . .                          | 90        |
| 9.7.2     | Knowledge Base Integration . . . . .                      | 91        |
| 9.7.3     | Verification Pipeline . . . . .                           | 91        |
| 9.8       | Educational Scaffolding Agents . . . . .                  | 92        |
| 9.8.1     | Mathematics Education . . . . .                           | 92        |
| 9.8.2     | Physics Education . . . . .                               | 93        |
| 9.9       | Research Agent Workflows . . . . .                        | 94        |
| 9.9.1     | Literature Review Agents . . . . .                        | 94        |
| 9.10      | Best Practices . . . . .                                  | 95        |
| 9.10.1    | Rigorous Verification . . . . .                           | 95        |
| 9.10.2    | Uncertainty Quantification . . . . .                      | 95        |
| 9.10.3    | Reproducibility . . . . .                                 | 96        |
| 9.10.4    | Domain Expert Collaboration . . . . .                     | 96        |
| 9.11      | Key Takeaways . . . . .                                   | 97        |
| <b>10</b> | <b>Common Failure Modes, Testing, and Fixes</b>           | <b>98</b> |
| 10.1      | Chapter Goals . . . . .                                   | 98        |
| 10.2      | Why Failures Are Different in Agentic Systems . . . . .   | 98        |
| 10.3      | Failure Taxonomy . . . . .                                | 98        |
| 10.3.1    | 1) Planning and Reasoning Failures . . . . .              | 98        |
| 10.3.2    | 2) Tooling and Integration Failures . . . . .             | 99        |
| 10.3.3    | 3) Context and Memory Failures . . . . .                  | 99        |
| 10.3.4    | 4) Safety and Policy Failures . . . . .                   | 99        |
| 10.3.5    | 5) Collaboration and Workflow Failures . . . . .          | 100       |
| 10.4      | Testing Strategy for Agentic Workflows . . . . .          | 100       |
| 10.5      | 1. Static and Structural Checks . . . . .                 | 100       |
| 10.6      | 2. Deterministic Unit Tests (Without LLM Calls) . . . . . | 100       |
| 10.7      | 3. Recorded Integration Tests (Golden Traces) . . . . .   | 101       |
| 10.8      | 4. Scenario and Adversarial Evaluations . . . . .         | 101       |
| 10.9      | 5. Production Guardrail Tests . . . . .                   | 101       |
| 10.10     | Practical Fix Patterns . . . . .                          | 101       |
| 10.10.1   | Pattern A: Contract Hardening . . . . .                   | 101       |
| 10.10.2   | Pattern B: Progressive Autonomy . . . . .                 | 102       |
| 10.10.3   | Pattern C: Two-Phase Execution . . . . .                  | 102       |
| 10.10.4   | Pattern D: Fallback and Circuit Breakers . . . . .        | 102       |
| 10.10.5   | Pattern E: Human-in-the-Loop Escalation . . . . .         | 102       |
| 10.11     | Incident Response Runbook (Template) . . . . .            | 102       |
| 10.12     | Metrics That Actually Matter . . . . .                    | 102       |
| 10.13     | Anti-Patterns to Avoid . . . . .                          | 102       |

|  |            |
|--|------------|
| 10.14A Minimal Reliability Checklist . . . . . | 103        |
| 10.15Chapter Summary . . . . .                 | 103        |
| <b>Bibliography</b>                            | <b>104</b> |

# Copyright

Copyright © 2026 Agentbook Contributors. All rights reserved.

This is an open-source manuscript distributed for educational purposes. No part of this work may be reproduced in any form or by any means without the prior written permission of the contributors, except for brief quotations in reviews.

# Preface

This book is for engineering leaders, platform teams, and practitioners who want to build reliable agentic workflows in real software systems. We focus on practical architecture, safe operations, and buildable examples rather than hype or speculative design.

You should be comfortable with Git, GitHub Actions, and basic software engineering practices. We assume familiarity with automation tools and CI/CD concepts.

# Conventions Used in This Book

- **Inline code** is shown in a monospaced font.
  - **File names** appear above code blocks (Example 4-2, etc.).
  - **Commands** use fenced code blocks with a language tag.
  - **Code samples** include a snippet-status label when version drift is likely:
    - **Runnable**: intended to run with minor adaptation, usually with a version/date note.
    - **Illustrative pseudocode**: architecture pattern, not copy/paste ready.
    - **Simplified**: runnable shape, but omits production details (auth, retries, error handling).
  - **Tip**, **Note**, and **Warning** callouts highlight critical guidance.
- Tip:** Prefer pinned action SHAs in CI workflows that handle secrets.
- Warning:** Always review agent-generated changes before merge, especially when credentials or production systems are involved.

# Chapter 1

## Introduction to Agentic Workflows

### 1.1 Chapter Preview

- Define agentic workflows and the roles they coordinate.
- Establish the terminology used throughout the book.
- Show where agentic workflows create leverage in practice.

### 1.2 What Are Agentic Workflows?

Agentic workflows represent a paradigm shift in how we approach software development and automation. Instead of writing explicit instructions for every task, we define goals and let AI agents determine the best path to achieve them.

#### 1.2.1 Key Concepts

**Agent:** An autonomous entity that can perceive its environment, make decisions, and take actions to achieve specific goals. In the context of software development, agents are AI-powered systems that can: - Read and understand code - Make modifications based on requirements - Test and verify changes - Interact with development tools and APIs

**Workflow:** A sequence of operations orchestrated to accomplish a complex task. Agentic workflows differ from traditional workflows by being: - **Adaptive:** Agents can modify their approach based on feedback - **Goal-oriented:** Focus on outcomes rather than rigid procedures - **Context-aware:** Understanding the broader context of their actions

#### 1.2.2 Terminology and Roles

To keep the manuscript consistent, we use the following terms throughout:

- **Agentic workflow** (primary term): A goal-directed, tool-using workflow executed by one or more agents.
- **Tool:** A capability exposed through a protocol (for example, an API, CLI, or MCP server).
- **Skill:** A packaged, reusable unit of instructions and/or code (see Skills and Tools Management).
- **Orchestrator:** The component that sequences work across agents.
- **Planner:** The component that decomposes goals into steps.
- **Executor:** The component that performs actions and records results.
- **Reviewer:** The component (often a human) that approves, rejects, or requests changes.

**Warning:** Prompt injection is a primary risk for agentic workflows. Treat external content as untrusted input and require explicit tool allowlists and human review for risky actions.

### 1.2.3 Why Agentic Workflows?

Traditional automation has limitations:

- **Rigid**: Predefined steps that can't adapt to unexpected situations
- **Fragile**: Breaking when conditions change
- **Limited scope**: Only handling well-defined, narrow tasks

Agentic workflows solve these problems by:

1. **Flexibility**: Adapting to changing requirements and conditions
2. **Intelligence**: Understanding intent and making informed decisions
3. **Scalability**: Handling increasingly complex tasks through composition

## 1.3 Real-World Applications

### 1.3.1 Software Development

- Automated code reviews and improvements
- Bug fixing and testing
- Documentation generation and updates
- Dependency management

### 1.3.2 Content Management

- Self-updating documentation (like this book!)
- Blog post generation and curation
- Translation and localization

### 1.3.3 Operations

- Infrastructure as Code management
- Automated incident response
- Performance optimization

## 1.4 The Agent Development Lifecycle

1. **Define Goals**: Specify what you want to achieve
2. **Configure Agents**: Set up agents with appropriate tools and permissions
3. **Execute Workflows**: Agents work toward goals autonomously
4. **Monitor and Refine**: Review outcomes and improve agent behavior
5. **Scale**: Compose multiple agents for complex tasks

## 1.5 Getting Started

To work with agentic workflows, you need:

- Understanding of AI/LLM capabilities and limitations
- Familiarity with the problem domain
- Tools and frameworks for agent development
- Infrastructure for agent execution

In the following chapters, we'll explore how to orchestrate agents, build scaffolding for agent-driven systems, and manage skills and tools effectively.

## 1.6 Key Takeaways

- Agentic workflows enable flexible, intelligent automation.
- Consistent terminology prevents confusion as systems scale.
- Security and human review guardrails are non-negotiable for production use.
- The rest of the book builds on these core concepts.

# Chapter 2

## Agent Orchestration

### 2.1 Chapter Preview

- Compare common orchestration patterns and when to use each.
- Map orchestration to roles (planner, executor, reviewer).
- Apply practical guardrails for coordination at scale.

### 2.2 Understanding Agent Orchestration

Agent orchestration is the art and science of coordinating multiple agents to work together toward common or complementary goals. Like conducting an orchestra where each musician plays their part, orchestration ensures agents collaborate effectively.

### 2.3 Orchestration Patterns

#### 2.3.1 Sequential Execution

Agents work one after another, each building on previous results.

```
Agent A -> Agent B -> Agent C -> Result
```

**Use cases:** - Code generation -> Testing -> Deployment - Data collection -> Analysis -> Reporting

#### 2.3.2 Parallel Execution

Multiple agents work simultaneously on independent tasks.

```
Agent A \
Agent B -> Aggregator -> Result
Agent C /
```

**Use cases:** - Multiple code reviews happening concurrently - Parallel data processing pipelines

#### 2.3.3 Hierarchical Execution

A supervisor agent delegates tasks to specialized worker agents.

```
Supervisor Agent
|--> Worker A
```

```
|--> Worker B  
`--> Worker C
```

**Use cases:** - Complex feature development with multiple components - Multi-stage testing and validation

### 2.3.4 Event-Driven Orchestration

Agents respond to events and trigger other agents.

```
Event -> Agent A -> Event -> Agent B -> Event -> Agent C
```

**Use cases:** - CI/CD pipelines - Automated issue management - Self-updating systems (like this book!)

## 2.4 Coordination Mechanisms

### 2.4.1 Message Passing

Agents communicate through messages containing: - Task descriptions - Context and data - Results and feedback

### 2.4.2 Shared State

Agents access common data stores: - Database - File system - Message queues - APIs

### 2.4.3 Direct Invocation

Agents directly call other agents: - Function calls - API requests - Workflow triggers

## 2.5 Best Practices

### 2.5.1 Clear Responsibilities

Define what each agent is responsible for:

```
agents:  
  code_reviewer:  
    role: Review code changes for quality and security  
    tools: [static_analysis, security_scanner]  
  
  test_runner:  
    role: Execute tests and report results  
    tools: [pytest, jest, test_framework]
```

### 2.5.2 Error Handling

Agents should handle failures gracefully: - Retry logic for transient failures - Fallback strategies - Clear error reporting - Rollback capabilities

### 2.5.3 Monitoring

Track agent performance: - Execution time - Success/failure rates - Resource usage - Output quality

## 2.5.4 Isolation

Keep agents independent:

- Minimize shared dependencies
- Use clear interfaces
- Version agent capabilities
- Test agents independently

**Note:** Orchestration should surface a clear audit trail: who decided, who executed, and who approved. Capture this early so later chapters can build on it.

## 2.6 Key Takeaways

- Orchestration is the coordination layer that turns goals into reliable execution.
- Use the simplest pattern that matches your dependency graph.
- Isolation, monitoring, and auditability are core design constraints.

## 2.7 Orchestration Frameworks

### 2.7.1 GitHub Actions

Workflow orchestration for GitHub repositories:

```
name: Agent Workflow
on: [push, pull_request]
jobs:
  review:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v6
      - name: Code Review Agent
        run: ./agents/review.sh
```

### 2.7.2 LangChain

LangChain (<https://python.langchain.com/>) is a Python framework for LLM applications:

**Snippet status:** Runnable example pattern (validated against LangChain v1 docs, Feb 2026; verify exact signature in your installed version).

```
from langchain.agents import create_agent

agent = create_agent(
    model="gpt-4o-mini",
    tools=tools,
    system_prompt="You are a helpful workflow orchestrator.",
)

result = agent.invoke({"messages": [{"role": "user", "content": "Your task here"}]})
```

**Note:** LangChain's agents API has evolved quickly. Prefer the current docs for exact signatures, and treat older snippets that pass `llm` directly as version-specific patterns.

### 2.7.3 Custom Orchestration

Build your own orchestrator:

```
class AgentOrchestrator:
    def __init__(self):
        self.agents = {}
```

```

def register(self, name, agent):
    self.agents[name] = agent

def execute_workflow(self, workflow_def):
    for step in workflow_def:
        agent = self.agents[step['agent']]
        result = agent.execute(step['task'])
        # Handle result and proceed

```

## 2.8 Real-World Example: Self-Updating Documentation

This book uses agent orchestration:

1. **Issue Monitor Agent:** Watches for new issues with suggestions
2. **Analysis Agent:** Determines if suggestion fits the book's scope
3. **Content Agent:** Writes or updates content
4. **Build Agent:** Generates markdown and PDF versions
5. **Publishing Agent:** Updates GitHub Pages
6. **Blog Agent:** Creates blog post about the update

All coordinated through GitHub Actions workflows!

## 2.9 Challenges and Solutions

### 2.9.1 Challenge: Agent Conflicts

**Solution:** Use locks, transactions, or coordinator patterns

### 2.9.2 Challenge: Debugging

**Solution:** Comprehensive logging, replay capabilities, visualization

### 2.9.3 Challenge: Performance

**Solution:** Caching, parallel execution, resource limits

### 2.9.4 Challenge: Versioning

**Solution:** Version agents and interfaces separately

## 2.10 Key Takeaways

- Orchestration coordinates multiple agents effectively
- Choose the right pattern for your use case
- Clear responsibilities and interfaces are essential
- Monitor and iterate on your orchestration strategies
- Use established frameworks when possible, but be ready to customize

# Chapter 3

## Agentic Scaffolding

### 3.1 Chapter Preview

- Identify the scaffolding layers that make agentic workflows reliable.
- Learn how to balance flexibility with safety controls.
- Map scaffolding decisions to operational risks.

### 3.2 What Is Agentic Scaffolding?

Agentic scaffolding is the infrastructure, frameworks, and patterns that enable agents to operate effectively. Just as scaffolding supports construction workers, agentic scaffolding provides the foundation for agent-driven development.

### 3.3 Core Components

#### 3.3.1 Tool Access Layer

Agents need controlled access to tools and APIs.

```
class ToolRegistry:  
    """Registry of tools available to agents"""  
  
    def __init__(self):  
        self._tools = {}  
  
    def register_tool(self, name, tool, permissions=None):  
        """Register a tool with optional permission constraints"""  
        self._tools[name] = {  
            'tool': tool,  
            'permissions': permissions or []  
        }  
  
    def get_tool(self, name, agent_id):  
        """Get tool if agent has permission"""  
        tool_config = self._tools.get(name)  
        if not tool_config:  
            raise ValueError(f"Tool {name} not found")  
  
        if self._check_permissions(agent_id, tool_config['permissions']):  
            return tool_config['tool']
```

```
raise PermissionError(f"Agent {agent_id} lacks permission for {name}")
```

### 3.3.2 Context Management

Maintain and share context between agent invocations.

```
class AgentContext:  
    """Manages context for agent execution"""  
  
    def __init__(self):  
        self.memory = {}  
        self.history = []  
  
    def store(self, key, value):  
        """Store information in context"""  
        self.memory[key] = value  
        self.history.append({  
            'action': 'store',  
            'key': key,  
            'timestamp': datetime.now()  
        })  
  
    def retrieve(self, key):  
        """Retrieve information from context"""  
        return self.memory.get(key)  
  
    def get_history(self):  
        """Get execution history"""  
        return self.history
```

### 3.3.3 Execution Environment

Provide safe, isolated environments for agent execution.

```
# Docker-based agent environment  
FROM python:3.11-slim  
  
# Install dependencies  
RUN pip install langchain openai requests  
  
# Set up workspace  
WORKDIR /agent_workspace  
  
# Security: Run as non-root user  
RUN useradd -m agent  
USER agent  
  
# Entry point for agent execution  
ENTRYPOINT ["python", "agent_runner.py"]
```

### 3.3.4 Communication Protocol

Standardize how agents communicate.

```
interface AgentMessage {  
    id: string;  
    sender: string;
```

```

recipient: string;
type: 'task' | 'result' | 'error' | 'query';
payload: any;
timestamp: Date;
metadata?: Record<string, any>;
}

class MessageBus {
    async send(message: AgentMessage): Promise<void> {
        // Route message to recipient
    }

    async subscribe(agentId: string, handler: MessageHandler): Promise<void> {
        // Subscribe agent to messages
    }
}

```

## 3.4 Scaffolding Patterns

### 3.4.1 Pattern 1: Tool Composition

Enable agents to combine tools effectively.

```

class ComposableTool:
    """Base class for composable tools"""

    def __init__(self, name, func, inputs, outputs):
        self.name = name
        self.func = func
        self.inputs = inputs
        self.outputs = outputs

    def compose_with(self, other_tool):
        """Compose this tool with another"""
        if self.outputs & other_tool.inputs:
            return CompositeTool([self, other_tool])
        raise ValueError("Tools cannot be composed - incompatible inputs/outputs")

    def execute(self, **kwargs):
        return self.func(**kwargs)

    # Usage
read_file = ComposableTool('read_file', read_func, set(), {'content'})
analyze_code = ComposableTool('analyze', analyze_func, {'content'}, {'issues'})
pipeline = read_file.compose_with(analyze_code)

```

### 3.4.2 Pattern 2: Skill Libraries

Organize reusable agent capabilities.

```

# skills/code_review.py
class CodeReviewSkill:
    """Skill for reviewing code changes"""

    def __init__(self, llm):
        self.llm = llm
        self.tools = ['git_diff', 'static_analysis', 'test_runner']

```

```

async def review_pull_request(self, pr_number):
    """Review a pull request"""
    diff = await self.get_diff(pr_number)
    issues = await self.analyze(diff)
    tests = await self.run_tests()
    return self.create_review(issues, tests)

# ... implementation details

# skills/_init_.py
from .code_review import CodeReviewSkill
from .documentation import DocumentationSkill
from .testing import TestingSkill

__all__ = ['CodeReviewSkill', 'DocumentationSkill', 'TestingSkill']

```

### 3.4.3 Pattern 3: Resource Management

Manage computational resources efficiently.

```

class ResourceManager:
    """Manages resources for agent execution"""

    def __init__(self, max_concurrent=5, timeout=300):
        self.max_concurrent = max_concurrent
        self.timeout = timeout
        self.active_agents = {}
        self.semaphore = asyncio.Semaphore(max_concurrent)

    async def execute_agent(self, agent_id, task):
        """Execute agent with resource limits"""
        async with self.semaphore:
            try:
                async with timeout(self.timeout):
                    result = await self._run_agent(agent_id, task)
                    return result
            except asyncio.TimeoutError:
                self._cleanup_agent(agent_id)
                raise AgentTimeoutError(f"Agent {agent_id} timed out")

```

### 3.4.4 Pattern 4: Observability

Monitor and debug agent behavior.

```

class AgentObserver:
    """Observes and logs agent behavior"""

    def __init__(self):
        self.logger = logging.getLogger('agent_observer')
        self.metrics = {}

    def log_execution(self, agent_id, task, result, duration):
        """Log agent execution"""
        self.logger.info(f"Agent {agent_id} executed {task} in {duration}s")
        self._update_metrics(agent_id, duration, result.success)

    def get_metrics(self, agent_id):

```

```

    """Get performance metrics"""
    return self.metrics.get(agent_id, {})

def export_trace(self, agent_id):
    """Export execution trace for debugging"""
    return self._build_trace(agent_id)

```

## 3.5 Building Scaffolding: Step by Step

### 3.5.1 Step 1: Define Your Agent Ecosystem

```

# agent_config.yaml
agents:
  content_writer:
    type: specialized
    tools: [markdown_editor, research_tool]
    max_execution_time: 600

  code_reviewer:
    type: specialized
    tools: [git, static_analyzer, test_runner]
    max_execution_time: 300

  orchestrator:
    type: coordinator
    tools: [task_queue, notification_service]
    manages: [content_writer, code_reviewer]

```

### 3.5.2 Step 2: Implement Tool Registry

Centralize tool access and management.

### 3.5.3 Step 3: Create Agent Templates

Provide starting points for common agent types.

```

# templates/base_agent.py
class BaseAgent(ABC):
    """Base template for all agents"""

    def __init__(self, agent_id, config):
        self.agent_id = agent_id
        self.config = config
        self.tools = self._load_tools()
        self.context = AgentContext()

    @abstractmethod
    async def execute(self, task):
        """Execute the agent's main task"""
        pass

    def _load_tools(self):
        """Load tools from registry"""
        return [get_tool(name) for name in self.config['tools']]

```

### 3.5.4 Step 4: Implement Error Recovery

Build resilience into your scaffolding.

```
class ResilientAgent:  
    """Agent with built-in error recovery"""\n\n    async def execute_with_recovery(self, task, max_retries=3):  
        """Execute with automatic retry on failure"""\n        for attempt in range(max_retries):\n            try:\n                result = await self.execute(task)\n                return result\n            except RecoverableError as e:\n                if attempt < max_retries - 1:\n                    await self._recover(e)\n                    continue\n                raise\n            except Exception as e:\n                self._log_error(e)\n                raise
```

**Warning:** Sandboxing and permission boundaries are not optional. Treat every tool invocation as a least-privilege request and validate all side effects in a separate review step.

## 3.6 Scaffolding for This Book

This book's scaffolding includes:

1. **GitHub Actions:** Workflow orchestration
2. **Issue Templates:** Structured input for suggestions
3. **Agent Scripts:** Python scripts for content management
4. **Tool Access:** Git, markdown processors, PDF generators
5. **State Management:** Git repository as persistent state
6. **Communication:** GitHub API for coordination

## 3.7 Best Practices

1. **Start Simple:** Build minimal scaffolding first, expand as needed
2. **Security First:** Implement permissions and isolation from the start
3. **Observability:** Log everything, you'll need it for debugging
4. **Version Control:** Version your scaffolding alongside your agents
5. **Documentation:** Document tools, APIs, and patterns clearly
6. **Testing:** Test your scaffolding independently of agents

## 3.8 Common Pitfalls

X **Over-engineering:** Don't build scaffolding for hypothetical needs X **Tight Coupling:** Keep agents loosely coupled to scaffolding X **Poor Error Handling:** Always plan for failure scenarios X **No Monitoring:** You can't improve what you can't measure X **Ignoring Security:** Security must be built in, not bolted on

## 3.9 Key Takeaways

- Scaffolding provides the foundation for effective agent operation

- Core components include tools, context, environment, and communication
- Patterns like composition and resource management improve scalability
- Build incrementally, focusing on security and observability
- Good scaffolding makes agents more capable and easier to manage

# Chapter 4

# Skills and Tools Management

## 4.1 Chapter Preview

- Define tools, skills, and how they map to operational workflows.
- Compare packaging formats and protocols for distributing skills.
- Walk through safe patterns for skill development and lifecycle management.

## 4.2 Understanding Skills vs. Tools

### 4.2.1 Tools

**Tools** are atomic capabilities that agents can use to interact with their environment. They are the building blocks of agent functionality.

Examples: - File system operations (read, write, delete) - API calls (GET, POST, PUT, DELETE) - Shell commands - Database queries

### 4.2.2 Skills

**Skills** are higher-level capabilities composed of multiple tools and logic. They represent complex behaviors that agents can learn and apply.

Examples: - Code review (using git, static analysis, test execution) - Documentation writing (using research, markdown editing, validation) - Bug fixing (using debugging, testing, code editing)

## 4.3 Tool Design Principles

### 4.3.1 Single Responsibility

Each tool should do one thing well.

```
# Good: Focused tool
class FileReader:
    """Reads content from files"""

    def read(self, filepath: str) -> str:
        with open(filepath, 'r') as f:
            return f.read()

# Bad: Tool doing too much
```

```

class FileManager:
    """Does everything with files"""

    def read(self, filepath): ...
    def write(self, filepath, content): ...
    def delete(self, filepath): ...
    def search(self, pattern): ...
    def backup(self, filepath): ...

```

### 4.3.2 Clear Interfaces

Tools should have well-defined inputs and outputs.

```

from typing import Protocol

class Tool(Protocol):
    """Interface for all tools"""

    name: str
    description: str

    def execute(self, **kwargs) -> dict:
        """Execute the tool with given parameters"""
        ...

    def get_schema(self) -> dict:
        """Get JSON schema for tool parameters"""
        ...

```

### 4.3.3 Error Handling

Tools must handle errors gracefully and provide useful feedback.

```

class WebScraperTool:
    """Tool for scraping web content"""

    def execute(self, url: str, timeout: int = 30) -> dict:
        try:
            response = requests.get(url, timeout=timeout)
            response.raise_for_status()
            return {
                'success': True,
                'content': response.text,
                'status_code': response.status_code
            }
        except requests.Timeout:
            return {
                'success': False,
                'error': 'Request timed out',
                'error_type': 'timeout'
            }
        except requests.RequestException as e:
            return {
                'success': False,
                'error': str(e),
                'error_type': 'request_error'
            }

```

#### 4.3.4 Documentation

Every tool needs clear documentation.

```
class GitDiffTool:  
    """  
    Tool for getting git diffs.  
  
    Capabilities:  
        - Get diff for specific files  
        - Get diff between commits  
        - Get diff for staged changes  
  
    Parameters:  
        filepath (str, optional): Specific file to diff  
        commit1 (str, optional): First commit hash  
        commit2 (str, optional): Second commit hash  
        staged (bool): Whether to show staged changes only  
  
    Returns:  
        dict: Contains 'diff' (str) and 'files_changed' (list)  
  
    Example:  
        >>> tool = GitDiffTool()  
        >>> result = tool.execute(staged=True)  
        >>> print(result['diff'])  
        """  
  
    def execute(self, **kwargs) -> dict:  
        # Implementation  
        pass
```

### 4.4 Creating Custom Tools

#### 4.4.1 Basic Tool Template

```
from typing import Any, Dict  
import json  
  
class CustomTool:  
    """Template for creating custom tools"""  
  
    def __init__(self, name: str, description: str):  
        self.name = name  
        self.description = description  
  
    def execute(self, **kwargs) -> Dict[str, Any]:  
        """  
        Execute the tool.  
  
        Override this method in your tool implementation.  
        """  
        raise NotImplementedError("Tool must implement execute method")  
  
    def validate_params(self, **kwargs) -> bool:  
        """  
        Validate tool parameters.
```

```

Override for custom validation logic.
"""
return True

def get_schema(self) -> Dict[str, Any]:
    """
    Return JSON schema for tool parameters.
    """
    return {
        'name': self.name,
        'description': self.description,
        'parameters': []
    }

```

#### 4.4.2 Example: Markdown Validation Tool

```

import re
from typing import Dict, Any, List

class MarkdownValidatorTool:
    """Validates markdown content for common issues"""

    def __init__(self):
        self.name = "markdown_validator"
        self.description = "Validates markdown files for common issues"

    def execute(self, content: str) -> Dict[str, Any]:
        """Validate markdown content"""
        issues = []

        # Check for broken links
        issues.extend(self._check_links(content))

        # Check for heading hierarchy
        issues.extend(self._check_headings(content))

        # Check for code block formatting
        issues.extend(self._check_code_blocks(content))

        return {
            'valid': len(issues) == 0,
            'issues': issues,
            'issue_count': len(issues)
        }

    def _check_links(self, content: str) -> List[Dict]:
        """Check for broken or malformed links"""
        issues = []
        links = re.findall(r'\[(\[^ ]+)\]\(((\^\))+)\\', content)

        for text, url in links:
            if not url:
                issues.append({
                    'type': 'broken_link',
                    'message': f'Empty URL in link: [{text}]()',
                    'severity': 'error'
                })

```

```

        return issues

def _check_headings(self, content: str) -> List[Dict]:
    """Check heading hierarchy"""
    issues = []
    lines = content.split('\n')
    prev_level = 0

    for i, line in enumerate(lines):
        if line.startswith('#'):
            level = len(line) - len(line.lstrip('#'))
            if level > prev_level + 1:
                issues.append({
                    'type': 'heading_skip',
                    'message': f'Heading level skipped at line {i+1}',
                    'severity': 'warning'
                })
            prev_level = level

    return issues

def _check_code_blocks(self, content: str) -> List[Dict]:
    """Check code block formatting"""
    issues = []
    backticks = re.findall(r'```, content)

    if len(backticks) % 2 != 0:
        issues.append({
            'type': 'unclosed_code_block',
            'message': 'Unclosed code block detected',
            'severity': 'error'
        })

    return issues

```

## 4.5 Agent Skills Standard (Primary Reference)

For practical interoperability, treat **Agent Skills** as the primary standard today. The authoritative docs are:

- Overview and motivation: <https://agentskills.io/home>
- Core concept page: <https://agentskills.io/what-are-skills>
- Full format specification: <https://agentskills.io/specification>
- Integration guidance: <https://agentskills.io/integrate-skills>

The current ecosystem signal is strongest around this filesystem-first model: a `SKILL.md` contract with progressive disclosure, plus optional `scripts/`, `references/`, and `assets/` directories.

**Placement note:** For OpenAI Codex auto-discovery, store repository skills under `.agents/skills/` (or user-level `~/.codex/skills/`). A plain top-level `skills/` folder is a useful wrapper convention but is not auto-discovered by default without additional wiring.

### 4.5.1 Canonical Layout

Example 4-1. `..agents/skills/code-review/`

```
.agents/
  skills/
```

```

code-review/
  SKILL.md
  manifest.json
  scripts/
    review.py
  references/
    rubric.md
  assets/
    example-diff.txt

```

Example 4-2. `agents/skills/code-review/SKILL.md`

```

---
name: code-review
description: Review pull requests for security, correctness, and clarity.
compatibility: Requires git and Python 3.11+
allowed-tools: Bash(git:*) Read
metadata:
  author: engineering-platform
  version: "1.2"
---

# Code Review Skill

## When to use
Use this skill when reviewing pull requests for correctness, security, and clarity.

## Workflow
1. Run `scripts/review.py --pr <number>`.
2. If policy checks fail, consult `references/rubric.md`.
3. Return findings grouped by severity and file.

```

**Tip:** Keep `SKILL.md` concise and front-load decision-critical instructions. Put deep references in `references/` and executable logic in `scripts/` so agents load content only when needed.

#### 4.5.2 Conformance Notes: Agent Skills vs. JSON-RPC Runtime Specs

The main discrepancy you will see across docs in the wild is **where standardization happens**:

- **Agent Skills** standardizes the **artifact format** (directory + `SKILL.md` schema + optional folders).
- Some alternative specs standardize a **remote runtime API** (often JSON-RPC-style methods such as `list`, `describe`, `execute`).

In production, the Agent Skills packaging approach currently has clearer multi-tool adoption because it works in both:

1. **Filesystem-based agents** (agent can `cat` and run local scripts).
2. **Tool-based agents** (host platform loads and mediates skill content).

JSON-RPC itself is battle-tested in other ecosystems (for example, Language Server Protocol, Ethereum node APIs, and MCP transport patterns), but there are still fewer public, concrete references to large-scale deployments of a dedicated JSON-RPC **skills runtime** than to plain `SKILL.md`-based workflows. For most teams, this makes Agent Skills the safest default and JSON-RPC skill runtimes an optional layering.

#### 4.5.3 Relationship to MCP

Use **Agent Skills** to define and distribute reusable capability packages. Use **MCP** to expose tools, data sources, or execution surfaces to models. In mature systems, these combine naturally: Agent Skills provide instructions and assets, while MCP provides controlled runtime tool access.

## 4.6 Skill Development

### 4.6.1 Skill Architecture

```
from typing import List, Dict, Any
from abc import ABC, abstractmethod

class Skill(ABC):
    """Base class for agent skills"""

    def __init__(self, name: str, tools: List[Tool]):
        self.name = name
        self.tools = {tool.name: tool for tool in tools}

    @abstractmethod
    async def execute(self, task: Dict[str, Any]) -> Dict[str, Any]:
        """Execute the skill"""
        pass

    def get_tool(self, name: str) -> Tool:
        """Get a tool by name"""
        return self.tools[name]

    def has_tool(self, name: str) -> bool:
        """Check if skill has a tool"""
        return name in self.tools
```

### 4.6.2 Example: Code Review Skill

```
class CodeReviewSkill(Skill):
    """Skill for reviewing code changes"""

    def __init__(self, llm, tools: List[Tool]):
        super().__init__("code_review", tools)
        self.llm = llm

    async def execute(self, task: Dict[str, Any]) -> Dict[str, Any]:
        """
        Execute code review.

        Task should contain:
        - pr_number: Pull request number
        - focus_areas: List of areas to focus on (optional)
        """
        pr_number = task['pr_number']
        focus_areas = task.get('focus_areas', ['bugs', 'security', 'performance'])

        # Step 1: Get code changes
        git_diff = self.get_tool('git_diff')
        diff_result = git_diff.execute(pr_number=pr_number)

        if not diff_result['success']:
            return {'success': False, 'error': 'Failed to get diff'}

        # Step 2: Run static analysis
        analyzer = self.get_tool('static_analyzer')
        analysis_result = analyzer.execute(
            diff=diff_result['diff'],
```

```

        focus=focus_areas
    )

    # Step 3: Run tests
    test_runner = self.get_tool('test_runner')
    test_result = test_runner.execute()

    # Step 4: Generate review using LLM
    review = await self._generate_review(
        diff_result['diff'],
        analysis_result['issues'],
        test_result
    )

    return {
        'success': True,
        'review': review,
        'static_analysis': analysis_result,
        'test_results': test_result
    }

async def _generate_review(self, diff, issues, tests):
    """Generate review using LLM"""
    prompt = f"""
Review the following code changes:

{diff}

Static analysis found these issues:
{json.dumps(issues, indent=2)}

Test results:
{json.dumps(tests, indent=2)}

Provide a comprehensive code review.
"""

    return await self.llm.generate(prompt)

```

## 4.7 Importing and Using Skills

### 4.7.1 Skill Registry

```

class SkillRegistry:
    """Central registry for skills"""

    def __init__(self):
        self._skills = {}

    def register(self, skill: Skill):
        """Register a skill"""
        self._skills[skill.name] = skill

    def get(self, name: str) -> Skill:
        """Get a skill by name"""
        if name not in self._skills:
            raise ValueError(f"Skill '{name}' not found")

```

```

        return self._skills[name]

    def list_skills(self) -> List[str]:
        """List all registered skills"""
        return list(self._skills.keys())

    def import_skill(self, module_path: str, skill_class: str):
        """Dynamically import and register a skill"""
        import importlib

        module = importlib.import_module(module_path)
        SkillClass = getattr(module, skill_class)

        # Instantiate and register
        skill = SkillClass()
        self.register(skill)

    # Usage
    registry = SkillRegistry()

    # Register built-in skills
    registry.register(CodeReviewSkill(llm, tools))
    registry.register(DocumentationSkill(llm, tools))

    # Import external skill
    registry.import_skill('external_skills.testing', 'TestGenerationSkill')

    # Use a skill
    code_review = registry.get('code_review')
    result = await code_review.execute({'pr_number': 123})

```

#### 4.7.2 Skill Composition

```

class CompositeSkill(Skill):
    """Skill composed of multiple sub-skills"""

    def __init__(self, name: str, skills: List[Skill]):
        self.name = name
        self.skills = {skill.name: skill for skill in skills}

        # Aggregate tools from all skills
        all_tools = []
        for skill in skills:
            all_tools.extend(skill.tools.values())

        super().__init__(name, list(set(all_tools)))

    async def execute(self, task: Dict[str, Any]) -> Dict[str, Any]:
        """Execute composed skill"""
        results = {}

        for skill_name, skill in self.skills.items():
            result = await skill.execute(task)
            results[skill_name] = result

        return {
            'success': all(r.get('success', False) for r in results.values()),
            'results': results
        }

```

```

    }

# Create composite skill
full_review = CompositeSkill('full_review', [
    CodeReviewSkill(llm, tools),
    SecurityAuditSkill(llm, tools),
    PerformanceAnalysisSkill(llm, tools)
])

```

## 4.8 Tool Discovery and Documentation

### 4.8.1 Self-Documenting Tools

```

class DocumentedTool:
    """Tool with built-in documentation"""

    def __init__(self):
        self.name = "example_tool"
        self.description = "Example tool with documentation"
        self.parameters = {
            'required': ['param1'],
            'optional': ['param2', 'param3'],
            'schema': {
                'param1': {'type': 'string', 'description': 'Required parameter'},
                'param2': {'type': 'int', 'description': 'Optional parameter'},
                'param3': {'type': 'bool', 'description': 'Flag parameter'}
            }
        }
        self.examples = [
            {
                'input': {'param1': 'value'},
                'output': {'success': True, 'result': 'output'}
            }
        ]

    def get_documentation(self) -> str:
        """Generate documentation for this tool"""
        doc = f"# {self.name}\n\n"
        doc += f"{self.description}\n\n"
        doc += "## Parameters\n\n"

        for param, schema in self.parameters['schema'].items():
            required = "Required" if param in self.parameters['required'] else "Optional"
            doc += f"- `{param}` ({schema['type']}), {required}: {schema['description']}\n"

        doc += "\n## Examples\n\n"
        for i, example in enumerate(self.examples, 1):
            doc += f"### Example {i}\n"
            doc += f"Input: `json.dumps(example['input'])`\n"
            doc += f"Output: `json.dumps(example['output'])`\n\n"

        return doc

```

## 4.9 Integrations: Connecting Tools to Real-World Surfaces

**Integrations** sit above tools and skills. They represent packaged connectors to real systems (chat apps, device surfaces, data sources, or automation backends) that deliver a coherent user experience. Think of them as the **distribution layer** for tools and skills: they bundle auth, event routing, permissions, and UX entry points.

**How integrations relate to tools and skills:**

- Tools are atomic actions (send a message, fetch a calendar event, post to Slack).
- Skills orchestrate tools to solve tasks (triage inbox, compile meeting notes, run a daily report).
- **Integrations** wrap tools + skills into deployable connectors with lifecycle management (pairing, secrets, rate limits, onboarding, and UI hooks).

In practice, a single integration might expose multiple tools and enable multiple skills. The integration is the bridge between agent capabilities and the messy realities of authentication, permissions, and channel-specific constraints.

## 4.10 Case Study: OpenClaw and pi-mono

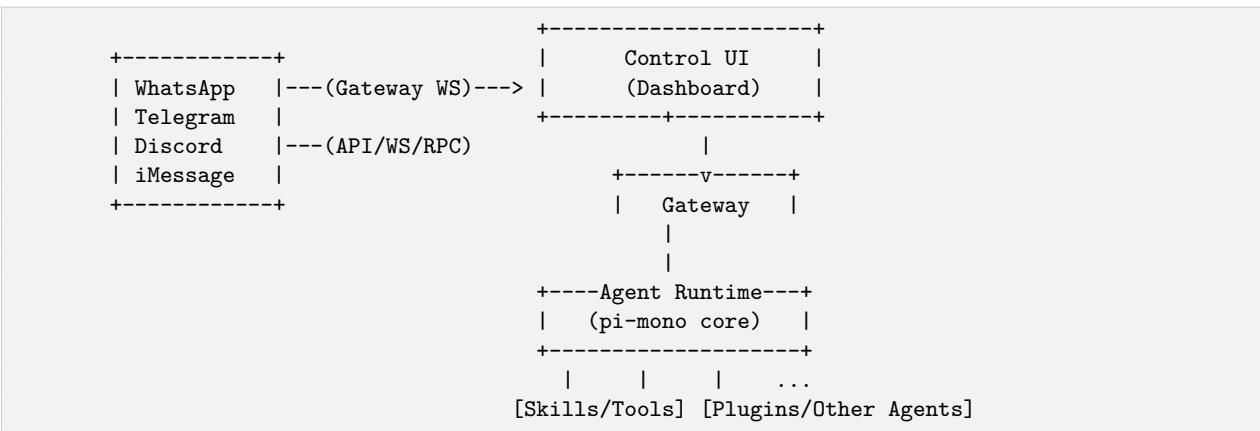
**Note:** OpenClaw and pi-mono are hypothetical composite examples used to illustrate architecture patterns. They are not official products as of 2026-02.

**OpenClaw** is a personal, local-first AI assistant that runs a gateway control plane and connects to many chat providers and device surfaces. It emphasizes multi-channel inboxes, tool access, and skill management inside a user-owned runtime.

OpenClaw is built on the **pi-mono** ecosystem. The pi-mono monorepo provides an agent runtime, tool calling infrastructure, and multi-provider LLM APIs that OpenClaw can leverage to keep the assistant portable across models and deployments.

### 4.10.1 OpenClaw Architecture in Detail

OpenClaw's architecture consists of several interconnected components:



**1. Gateway Control Plane** - Central hub orchestrating all user input/output and messaging channels  
- Exposes a WebSocket server (default: `ws://127.0.0.1:18789`)  
- Handles session state, permissions, and authentication  
- Supports local and mesh/LAN deployment via Tailscale (<https://tailscale.com/>) or similar

**2. Pi Agent Runtime (pi-mono)** - Core single-agent execution environment  
- Maintains long-lived agent state, memory, skills, and tool access  
- Handles multi-turn conversation, contextual memory, and tool/plugin invocation  
- Orchestrates external API/model calls (OpenAI, Anthropic, local models via Ollama <https://ollama.com/>)  
- Persistent storage (SQLite, Postgres, Redis) for memory and context

**3. Multi-Agent Framework** - Support for swarms of specialized agents (“nodes”) handling domain-specific automations - Agents coordinate via shared memory and routing protocols managed by the Gateway - Each agent can be sandboxed (Docker/isolation) for security - Developers build custom agents via TypeScript/YAML plugins

**4. Extensible Skills/Plugin Ecosystem** - Skills expand the agent’s abilities: file automation, web scraping, email, calendar - Plugins are hot-reloadable and built in TypeScript - Community skill marketplace for sharing and discovery

#### 4.10.2 Key Design Principles

1. **Privacy-First:** All state and memory default to local storage—data never leaves the device unless explicitly configured
2. **BYOM (Bring Your Own Model):** Seamlessly supports cloud LLMs and local inference
3. **Proactive Behavior:** “Heartbeat” feature enables autonomous wake-up and environment monitoring
4. **Persistent Memory:** Learns and adapts over long-term interactions

Key takeaways for skills/tools architecture:

- **Gateway + runtime separation** keeps tools and skills consistent while integrations change: the gateway handles channels and routing, while pi-mono-style runtimes handle tool execution.
- **Integration catalogs** (like OpenClaw’s integrations list and skill registry) are a user-facing map of capability. They surface what tools can do and what skills are available without forcing users to understand low-level APIs.
- **Skills become reusable assets** once tied to integrations: a “Slack triage” skill can target different workspaces without changing the underlying tools, as long as the integration provides the same tool contracts.

### 4.11 Related Architectures and Frameworks

Several other frameworks share architectural patterns with OpenClaw:

#### 4.11.1 LangChain and LangGraph

LangChain (<https://python.langchain.com/>) provides composable building blocks for LLM applications.

**Snippet status:** Illustrative pseudocode (API names vary across LangChain versions).

```
from langchain.agents import create_agent
from langchain_core.tools import tool

@tool
def search_documentation(query: str) -> str:
    """Search project documentation for relevant information."""
    # Implementation
    return "..."

agent = create_agent(
    model="gpt-4o-mini",
    tools=[search_documentation],
    system_prompt="Use tools when needed, then summarize clearly.",
)
result = agent.invoke({"messages": [{"role": "user", "content": "Find testing docs"}]})
```

**Shared patterns with OpenClaw:** Tool registration, agent composition, memory management.

LangGraph (<https://langchain-ai.github.io/langgraph/>) extends LangChain with graph-based agent orchestration.

#### 4.11.2 CrewAI

CrewAI (<https://docs.crewai.com/>) focuses on multi-agent collaboration with role-based specialization:

```
from crewai import Agent, Task, Crew

researcher = Agent(
    role='Senior Researcher',
    goal='Discover new insights',
    backstory='Expert in finding and analyzing information',
    tools=[search_tool, analysis_tool]
)

writer = Agent(
    role='Technical Writer',
    goal='Create clear documentation',
    backstory='Skilled at explaining complex topics',
    tools=[writing_tool]
)

crew = Crew(
    agents=[researcher, writer],
    tasks=[research_task, writing_task],
    process=Process.sequential
)
```

**Shared patterns with OpenClaw:** Role-based agents, sequential and parallel execution, tool assignment per agent.

#### 4.11.3 Microsoft Semantic Kernel

Semantic Kernel (<https://learn.microsoft.com/semantic-kernel/>) emphasizes enterprise integration and plugin architecture:

```
var kernel = Kernel.CreateBuilder()
    .AddOpenAIChatCompletion("gpt-4", apiKey)
    .Build();

// Import plugins
kernel.ImportPluginFromType<TimePlugin>();
kernel.ImportPluginFromType<FileIOPPlugin>();

// Create agent with plugins
var agent = new ChatCompletionAgent {
    Kernel = kernel,
    Name = "ProjectAssistant",
    Instructions = "Help manage project tasks and documentation"
};
```

**Shared patterns with OpenClaw:** Plugin system, kernel/runtime separation, enterprise-ready design.

#### 4.11.4 AutoGen

AutoGen (<https://microsoft.github.io/autogen/>) specializes in conversational multi-agent systems:

```
from autogen import AssistantAgent, UserProxyAgent

assistant = AssistantAgent(
    name="coding_assistant",
```

```

        llm_config={"model": "gpt-4",
                     system_message="You are a helpful coding assistant."
                 }

        user_proxy = UserProxyAgent(
            name="user",
            human_input_mode="NEVER",
            code_execution_config={"work_dir": "coding"}
        )

    # Agents collaborate through conversation
    user_proxy.initiate_chat(assistant, message="Create a Python web scraper")

```

**Shared patterns with OpenClaw:** Agent-to-agent communication, code execution environments, conversation-driven workflows.

#### 4.11.5 Comparing Architecture Patterns

| Feature                        | OpenClaw            | LangChain        | CrewAI             | Semantic Kernel    | AutoGen                  |
|--------------------------------|---------------------|------------------|--------------------|--------------------|--------------------------|
| <b>Primary Focus</b>           | Personal assistant  | LLM app building | Team collaboration | Enterprise plugins | Multi-agent chat         |
| <b>Runtime</b>                 | Local-first         | Flexible         | Python process     | .NET/Python        | Python process           |
| <b>Multi-Agent Tool System</b> | Via swarms          | Via LangGraph    | Built-in           | Via agents         | Built-in                 |
| <b>Memory</b>                  | Plugin-based        | Tool decorators  | Tool assignment    | Plugin imports     | Tool decorators          |
| <b>Best For</b>                | Persistent local    | Configurable     | Per-agent          | Session-based      | Conversation             |
|                                | Personal automation | Prototyping      | Complex workflows  | Enterprise apps    | Research/experimentation |

## 4.12 MCP: Modern Tooling and Adoption

The **Model Context Protocol (MCP)** (<https://modelcontextprotocol.io/>) has become a practical standard for connecting agents to tools and data sources. Today, MCP is less about novel capability and more about **reliable interoperability**: the same tool server can be used by multiple agent clients with consistent schemas, permissions, and response formats.

#### 4.12.1 What MCP Brings to Tools

- **Portable tool definitions:** JSON schemas and well-known server metadata make tools discoverable across clients.
- **Safer tool execution:** capability-scoped permissions, explicit parameters, and auditable tool calls.
- **Composable context:** servers can enrich model context with structured resources (files, APIs, or databases) without bespoke glue code.

#### 4.12.2 Common Usage Patterns

1. **Server-based tool catalogs**
  - Teams deploy MCP servers per domain (e.g., “repo-tools”, “ops-tools”, “research-tools”).
  - Clients discover available tools at runtime and choose based on metadata, not hardcoded lists.
2. **Context stitching**

- Agents gather context from multiple servers (docs, tickets, metrics) and assemble it into a task-specific prompt.
- The server provides structured resources so the client can keep the prompt lean.

### 3. Permission-first workflows

- Tool calls are scoped by project, environment, or role.
- Audit logs track who called what tool with which inputs.

### 4. Fallback-first reliability

- Clients maintain fallbacks when a server is down (cached data, read-only mirrors, or alternative tool servers).

## 4.12.3 Acceptance Across Major Clients

MCP is broadly accepted as a **tooling interoperability layer**. The specifics vary by vendor, but the pattern is consistent: MCP servers expose the tools and resources, while clients orchestrate tool calls and manage safety policies.

- **Codex (GPT-5.2-Codex)** (<https://openai.com/index/introducing-codex/>)  
Codex clients commonly use MCP servers to standardize tool access (repo browsing, test execution, task automation). Codex also supports skills packaged with `SKILL.md` and progressive disclosure (see <https://platform.openai.com/docs/guides/codex/skills>). The main adoption pattern is organization-level MCP servers that provide consistent tools across multiple repos.
- **GitHub Copilot** (<https://docs.github.com/en/copilot>)  
Copilot deployments increasingly treat MCP as a bridge between editor experiences and organization tooling. This typically means MCP servers that expose repo-aware tools (search, CI status, documentation retrieval) so the assistant can operate with consistent, policy-driven access.
- **Claude** (<https://code.claude.com/docs>)  
Claude integrations often use MCP to provide structured context sources (knowledge bases, issue trackers, dashboards). The MCP server becomes the policy boundary, while the client focuses on prompt composition and response quality.

## 4.12.4 Practical Guidance for Authors and Teams

- **Document your MCP servers** like any other tool: include schemas, permissions, and usage examples.
- **Version tool contracts** so clients can adopt changes incrementally.
- **Prefer narrow, composable tools** over large monolithic endpoints.
- **Treat MCP as infrastructure**: invest in uptime, monitoring, and security reviews.

## 4.13 Best Practices

### 4.13.1 Version Tools and Skills

```
class VersionedTool:
    def __init__(self, version: str):
        self.version = version
        self.name = f"{self.__class__.__name__}_v{version}"
```

### 4.13.2 Test Independently

```
# test_tools.py
import pytest

def test_markdown_validator():
```

```

tool = MarkdownValidatorTool()

# Test valid markdown
valid_md = "# Header\n\nContent"
result = tool.execute(valid_md)
assert result['valid']

# Test invalid markdown
invalid_md = "```python\ncode without closing"
result = tool.execute(invalid_md)
assert not result['valid']
assert any(i['type'] == 'unclosed_code_block' for i in result['issues'])

```

#### 4.13.3 Provide Fallbacks

```

class ResilientTool:
    def __init__(self, primary_impl, fallback_impl):
        self.primary = primary_impl
        self.fallback = fallback_impl

    def execute(self, **kwargs):
        try:
            return self.primary.execute(**kwargs)
        except Exception as e:
            logger.warning(f"Primary implementation failed: {e}")
            return self.fallback.execute(**kwargs)

```

#### 4.13.4 Monitor Usage

```

class MonitoredTool:
    def __init__(self, tool, metrics_collector):
        self.tool = tool
        self.metrics = metrics_collector

    def execute(self, **kwargs):
        start = time.time()
        try:
            result = self.tool.execute(**kwargs)
            self.metrics.record_success(self.tool.name, time.time() - start)
            return result
        except Exception as e:
            self.metrics.record_failure(self.tool.name, str(e))
            raise

```

### 4.14 Emerging Standards: AGENTS.md

#### 4.14.1 The AGENTS.md Pseudo-Standard

**AGENTS.md** has emerged as an open pseudo-standard for providing AI coding agents with project-specific instructions. Think of it as a “README for agents”—offering structured, machine-readable guidance that helps agents understand how to work within a codebase.

##### 4.14.1.1 Purpose and Benefits

- **Consistent Instructions:** All agents receive the same project-specific guidance

- **Rapid Onboarding:** New agent sessions understand the project immediately
- **Safety Boundaries:** Clear boundaries prevent accidental damage to protected files
- **Maintainability:** Single source of truth for agent behavior in a project

#### 4.14.1.2 Structure and Placement

AGENTS.md files can be placed hierarchically in a project:

```
project/
|-- AGENTS.md          # Root-level instructions (project-wide)
|-- src/
|   '-- AGENTS.md      # Module-specific instructions
|-- tests/
|   '-- AGENTS.md      # Testing conventions
`-- docs/
    '-- AGENTS.md      # Documentation guidelines
```

Agents use the nearest AGENTS.md file, enabling scoped configuration for monorepos or complex projects.

#### 4.14.1.3 Example AGENTS.md

```
# AGENTS.md

## Project Overview
This is a TypeScript web application using Express.js and React.

## Setup Instructions
npm install
npm run dev

## Coding Conventions
- Language: TypeScript 5.x
- Style guide: Airbnb
- Formatting: Prettier with provided config
- Test framework: Jest

## Build and Deploy
- Build: `npm run build`
- Test: `npm test`
- Deploy: CI/CD via GitHub Actions

## Agent-Specific Notes
- Always run `npm run lint` before committing
- Never modify files in `vendor/` or `/.github/workflows/`
- Secrets are in environment variables, never hardcoded
- All API endpoints require authentication middleware
```

#### 4.14.2 Related Standards: Skills and Capabilities

While AGENTS.md has achieved broad adoption as the standard for project-level agent instructions, the space continues to evolve. Several related concepts are under discussion in the community:

##### Skills Documentation

There is no formal skills.md standard, but skill documentation patterns are emerging:

- **Skill catalogs** listing available agent capabilities
- **Capability declarations** specifying what an agent can do

- **Dependency manifests** defining tool and skill requirements

### Personality and Values

Some frameworks experiment with “soul” or personality configuration. Note that “soul” is a metaphorical term used in some AI agent frameworks to describe an agent’s core personality, values, and behavioral guidelines—it’s industry jargon rather than a formal technical specification:

- **System prompts** defining agent persona and communication style
- **Value alignment** specifying ethical guidelines and constraints
- **Behavioral constraints** limiting what agents should and shouldn’t do

Currently, these are implemented in vendor-specific formats rather than open standards. The community continues to discuss whether formalization is needed.

## 4.15 How Agents Become Aware of Imports

One of the most practical challenges in agentic development is helping agents understand a codebase’s import structure and dependencies. When an agent modifies code, it must know what modules are available, where they come from, and how to properly reference them.

### 4.15.1 The Import Awareness Problem

When agents generate or modify code, they face several import-related challenges:

1. **Missing imports**: Adding code that uses undefined symbols
2. **Incorrect import paths**: Using wrong relative or absolute paths
3. **Circular dependencies**: Creating imports that cause circular reference errors
4. **Unused imports**: Leaving orphan imports after code changes
5. **Conflicting names**: Importing symbols that shadow existing names

### 4.15.2 Mechanisms for Import Discovery

Modern coding agents use multiple strategies to understand imports:

#### 4.15.2.1 Static Analysis Tools

Agents leverage language servers and static analyzers to understand import structure:

```
class ImportAnalyzer:
    """Analyze imports using static analysis"""

    def __init__(self, workspace_root: str):
        self.workspace = workspace_root
        self.import_graph = {}

    def analyze_file(self, filepath: str) -> dict:
        """Extract import information from a file"""
        with open(filepath) as f:
            content = f.read()

        # Parse AST to find imports
        tree = ast.parse(content)
        imports = []

        for node in ast.walk(tree):
            if isinstance(node, ast.Import):
                for alias in node.names:
                    imports.append({
```

```

        'type': 'import',
        'module': alias.name,
        'alias': alias.asname
    })
elif isinstance(node, ast.ImportFrom):
    imports.append({
        'type': 'from_import',
        'module': node.module,
        'names': [a.name for a in node.names],
        'level': node.level # relative import level
    })

return {
    'file': filepath,
    'imports': imports,
    'defined_symbols': self._extract_definitions(tree)
}

def build_dependency_graph(self) -> dict:
    """Build a graph of all file dependencies"""
    for filepath in self._find_source_files():
        analysis = self.analyze_file(filepath)
        self.import_graph[filepath] = analysis
    return self.import_graph

```

#### 4.15.2.2 Language Server Protocol (LSP)

Language servers provide real-time import information that agents can query:

```

class LSPIimportProvider:
    """Use LSP to discover available imports"""

    async def get_import_suggestions(self, symbol: str, context_file: str) -> list:
        """Get import suggestions for an undefined symbol"""

        # Query language server for symbol locations
        response = await self.lsp_client.request('textDocument/codeAction', {
            'textDocument': {'uri': context_file},
            'context': {
                'diagnostics': [
                    'message': f"Cannot find name '{symbol}'"
                ]
            }
        })

        # Extract import suggestions from code actions
        suggestions = []
        for action in response:
            if 'import' in action.get('title', '').lower():
                suggestions.append({
                    'import_statement': action['edit']['changes'],
                    'source': action.get('title')
                })

        return suggestions

    async def get_exported_symbols(self, module_path: str) -> list:
        """Get all exported symbols from a module"""

```

```

# Use workspace/symbol to find exports
symbols = await self.lsp_client.request('workspace/symbol', {
    'query': '',
    'uri': module_path
})

return [s['name'] for s in symbols if s.get('kind') in EXPORTABLE_KINDS]

```

#### 4.15.2.3 Project Configuration Files

Agents read configuration files to understand module resolution:

```

class ProjectConfigReader:
    """Read project configs to understand import paths"""

    def get_import_config(self, project_root: str) -> dict:
        """Extract import configuration from project files"""

        config = {
            'base_paths': [],
            'aliases': {},
            'external_packages': []
        }

        # TypeScript/JavaScript: tsconfig.json, jsconfig.json
        tsconfig_path = os.path.join(project_root, 'tsconfig.json')
        if os.path.exists(tsconfig_path):
            with open(tsconfig_path) as f:
                tsconfig = json.load(f)

            compiler_opts = tsconfig.get('compilerOptions', {})
            config['base_paths'].append(compiler_opts.get('baseUrl', '.'))
            config['aliases'] = compiler_opts.get('paths', {})

        # Python: pyproject.toml, setup.py
        pyproject_path = os.path.join(project_root, 'pyproject.toml')
        if os.path.exists(pyproject_path):
            with open(pyproject_path) as f:
                pyproject = toml.load(f)

        # Extract package paths from tool.setuptools or poetry config
        if 'tool' in pyproject:
            if 'setuptools' in pyproject['tool']:
                config['base_paths'].extend(
                    pyproject['tool']['setuptools'].get('package-dir', {}).values()
                )

        return config

```

#### 4.15.2.4 Package Manifest Analysis

Agents check package manifests to know what's available:

```

class PackageManifestReader:
    """Read package manifests to understand available dependencies"""

    def get_available_packages(self, project_root: str) -> dict:
        """Get list of available packages from manifest"""

```

```

packages = {'direct': [], 'transitive': []}

# Node.js: package.json
package_json = os.path.join(project_root, 'package.json')
if os.path.exists(package_json):
    with open(package_json) as f:
        pkg = json.load(f)
    packages['direct'].extend(pkg.get('dependencies', {}).keys())
    packages['direct'].extend(pkg.get('devDependencies', {}).keys())

# Python: requirements.txt, Pipfile, pyproject.toml
requirements = os.path.join(project_root, 'requirements.txt')
if os.path.exists(requirements):
    with open(requirements) as f:
        for line in f:
            line = line.strip()
            if line and not line.startswith('#'):
                # Extract package name (before version specifier)
                pkg_name = re.split(r'[^>=!=]', line)[0].strip()
                packages['direct'].append(pkg_name)

return packages

```

### 4.15.3 Best Practices for Import-Aware Agents

#### 4.15.3.1 Document Import Conventions in AGENTS.md

For standardized terminology (artefact, discovery, import, install, activate) and trust boundaries, see Discovery and Imports. In this section we apply those concepts to codebase-level import resolution.

Include import guidance in your project's AGENTS.md:

```

## Import Conventions

### Path Resolution
- Use absolute imports from `src/` as the base
- Prefer named exports over default exports
- Group imports: stdlib, external packages, local modules

### Example Import Order
```python
# Standard library
import os
import sys
from typing import Dict, List

# Third-party packages
import requests
from pydantic import BaseModel

# Local modules
from src.utils import helpers
from src.models import User
```

```

### 4.15.4 Alias Conventions

- @/ maps to src/

- @components/ maps to src/components/

```
#### Use Import Auto-Fix Tools

Configure agents to use automatic import fixers:

```python
class ImportAutoFixer:
    """Automatically fix import issues in agent-generated code"""

    def __init__(self, tools: List[Tool]):
        self.isort = tools.get('isort') # Python import sorting
        self.eslint = tools.get('eslint') # JS/TS import fixing

    async def fix_imports(self, filepath: str) -> dict:
        """Fix and organize imports in a file"""

        results = {'fixed': [], 'errors': []}

        if filepath.endswith('.py'):
            # Run isort for Python
            result = await self.isort.execute(filepath)
            if result['success']:
                results['fixed'].append('isort: organized imports')

            # Run autoflake to remove unused imports
            result = await self.autoflake.execute(
                filepath,
                remove_unused_imports=True
            )
            if result['success']:
                results['fixed'].append('autoflake: removed unused')

        elif filepath.endswith('.ts', '.tsx', '.js', '.jsx'):
            # Run eslint with import rules
            result = await self.eslint.execute(
                filepath,
                fix=True,
                rules=['import/order', 'unused-imports/no-unused-imports']
            )
            if result['success']:
                results['fixed'].append('eslint: fixed imports')

        return results
```

```

#### 4.15.4.1 Validate Imports Before Committing

Add import validation to agent workflows:

```
# .github/workflows/validate-imports.yml
name: Validate Imports
on: [pull_request]

jobs:
  check-imports:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v6
```

```

- name: Check Python imports
  run: |
    pip install isort autoflake
    isort --check-only --diff .
    autoflake --check --remove-all-unused-imports -r .

- name: Check TypeScript imports
  run: |
    npm ci
    npx eslint --rule 'import/no-unresolved: error' .

```

#### 4.15.5 Import Awareness in Multi-Agent Systems

When multiple agents collaborate, maintaining consistent import awareness requires coordination:

```

class SharedImportContext:
    """Shared import context for multi-agent systems"""

    def __init__(self):
        self.import_cache = {}
        self.pending_additions = []

    def register_new_export(self, module: str, symbol: str, agent_id: str):
        """Register a new export created by an agent"""
        if module not in self.import_cache:
            self.import_cache[module] = []

        self.import_cache[module].append({
            'symbol': symbol,
            'added_by': agent_id,
            'timestamp': datetime.now()
        })

    def query_available_imports(self, symbol: str) -> List[dict]:
        """Query where a symbol can be imported from"""
        results = []
        for module, exports in self.import_cache.items():
            for export in exports:
                if export['symbol'] == symbol:
                    results.append({
                        'module': module,
                        'symbol': symbol,
                        'import_statement': f"from {module} import {symbol}"
                    })
        return results

```

Understanding how agents discover and manage imports is essential for building reliable agentic coding systems. The combination of static analysis, language servers, project configuration, and clear documentation ensures agents can write code that integrates correctly with existing codebases.

### 4.16 Key Takeaways

- Tools are atomic capabilities; skills are composed behaviors
- Design tools with single responsibility and clear interfaces
- Skills orchestrate multiple tools to accomplish complex tasks

- Use registries for discovery and management
- Skills can be composed to create more powerful capabilities
- Always document, test, and version your tools and skills
- Monitor usage to identify issues and optimization opportunities
- AGENTS.md is the emerging standard for project-level agent instructions
- Skills Protocol defines how runtimes execute skills; Agent Skills defines how they are packaged
- MCP standardizes tool interoperability across clients and hosts
- Import awareness requires combining static analysis, LSP, and project configuration
- OpenClaw, LangChain, CrewAI, and similar frameworks share common patterns for tool and skill management

# Chapter 5

## Discovery and Imports

### 5.1 Chapter Preview

This chapter standardizes language that often gets overloaded in agentic systems. We define what an artefact is, what discovery means in practice, and how import, install, and activate are separate operations with different safety controls.

- Build a taxonomy for what gets discovered: tools, skills, agents, and workflow fragments.
- Compare discovery mechanisms: local scan, registry, domain convention, and explicit configuration.
- Disambiguate import/install/activate and map each step to trust and supply-chain controls.

### 5.2 Terminology Baseline for the Rest of the Book

To keep later chapters precise, this chapter uses five core terms consistently:

- **artefact**: any reusable unit a workflow can reference (tool endpoint metadata, a skill bundle, an agent definition, or a workflow fragment).
- **discovery**: the process of finding candidate artefacts.
- **import**: bringing a discovered artefact into the current resolution/evaluation context.
- **install**: fetching and persisting an artefact (typically pinned), plus integrity metadata.
- **activate**: making an installed/imported artefact callable by an agent in a specific run context.

**Standardization rule:** Use these verbs literally. Do not use “import” when you mean “install,” and do not use “install” when you mean “activate.”

### 5.3 A Taxonomy That Disambiguates What We Are Discovering

#### 5.3.1 1) Tool artefacts

A **tool** is an executable capability exposed via a protocol or command surface.

- **Identity**: endpoint identity (for example, an MCP server URL + auth context).
- **Interface**: enumerated callable tools (names, schemas, permissions).

Discovery usually finds endpoints first; tool enumeration happens after connection.

#### 5.3.2 2) Skill artefacts

A **skill** is a packaged reusable bundle of instructions, templates, and optional scripts.

- **Identity**: bundle source (repo path, registry coordinate, version/digest).

- **Interface:** documented entrypoints, expected inputs/outputs, and policy constraints.

### 5.3.3 3) Agent artefacts

An **agent** artefact is a role/config definition (persona, constraints, and operating policy).

- **Identity:** a named definition file and version.
- **Interface:** responsibilities, boundaries, and allowed capability set.

### 5.3.4 4) Workflow-fragment artefacts

A **workflow fragment** is a reusable partial workflow (for example, a GH-AW component).

- **Identity:** source file path or import address.
- **Interface:** parameters, expected context, and emitted outputs.

### 5.3.5 Confusing cases to stop using

- **tool != skill:** tools execute capabilities; skills package guidance/assets.
- **skill != agent:** skills are reusable bundles; agents are operating roles.
- **agent != workflow fragment:** an agent is an actor; a fragment is orchestration structure.

When this book says “discover capabilities,” read it as “discover artefacts, then import/install/activate according to type.”

## 5.4 Discovery Mechanisms

Discovery is how runtimes gather candidate artefacts before selection.

### 5.4.1 Local scan

Scan repository paths and conventions (for example, `.github/workflows/`, `skills/`, `agents/`).

- **Pros:** low latency, high transparency, easy review in code.
- **Cons:** limited scope, convention drift in large monorepos.

### 5.4.2 Registry discovery

Query a curated index/marketplace for artefacts.

- **Pros:** centralized metadata, version visibility, governance hooks.
- **Cons:** trust shifts to registry policy; namespace collisions possible.

### 5.4.3 Domain-convention discovery

Resolve artefacts via domain naming conventions (for example, `.well-known`-style descriptors).

- **Pros:** interoperable discovery across organizational boundaries.
- **Cons:** conventions may be ecosystem-specific, not always standardized.

### 5.4.4 Explicit configuration

Use a pinned manifest that enumerates allowed sources and versions.

- **Pros:** strongest reproducibility and auditability.
- **Cons:** less flexible; requires deliberate updates.

**Decision rule:** if provenance cannot be authenticated, prefer explicit configuration over dynamic discovery.

## 5.5 Import, Install, Activate: Three Different Operations

### 5.5.1 Import

Import brings an artefact into the current resolution context.

- Language/module example: `from src.utils import helpers`
- GH-AW example: `imports: [shared/common-tools.md]`

### 5.5.2 Install

Install fetches and persists artefacts for repeatable use.

- Example: store `skill-x@1.4.2` with checksum/signature metadata.
- Example: lock workflow component revision to a commit digest.

### 5.5.3 Activate

Activate makes an artefact callable under policy.

- Example: expose only `bash` and `edit` tools to a CI agent.
- Example: enable a skill only after approval gate passes.

A practical sequence is often: `discover -> select -> import/install -> activate -> execute`.

## 5.6 Trust Boundaries and Supply Chain (Compact Model)

Each stage has distinct risks and controls:

- **Integrity** (was artefact tampered with?): checksums, signatures.
- **Authenticity** (who published it?): identity verification, trusted publisher lists.
- **Provenance** (how was it built?): attestations/SBOM, reproducible build metadata.
- **Capability safety** (what can it do?): least privilege, sandboxing, constrained outputs.

Control mapping:

- **Discovery controls**: allowlists of domains/registries.
- **Import/install controls**: pinning + checksum/signature verification.
- **Activation controls**: permission gates, scoped credentials, sandbox profiles.
- **Runtime controls**: audit logs, safe outputs, and policy evaluation traces.

## 5.7 Worked Examples

### 5.7.1 Example A: GH-AW workflow-fragment import

```
# .github/workflows/docs-refresh.md
name: Docs Refresh
on:
  workflow_dispatch:
imports:
  - shared/common-tools.md

permissions:
  contents: read
```

Interpretation: `- shared/common-tools.md` is a **workflow-fragment artefact**. `- imports` is the **import** operation. A separate policy decides whether imported tools are **activated**.

### 5.7.2 Example B: AGENTS.md import conventions as resolution policy

```
## Import Conventions
- Prefer absolute imports from `src/`  

- Group imports: stdlib, third-party, local  

- Use `@/` alias for `src/`
```

Interpretation: - This does **not** install dependencies. - It standardizes **import resolution behavior** so agents generate consistent code. - Activation still depends on tool/runtime permissions.

### 5.7.3 Example C: Capability discovery and activation policy

```
# policies/capability-sources.yml
allowed_domains:
  - "skills.example.com"
allowed_registries:
  - "registry.internal/agent-skills"
pinned:
  "registry.internal/agent-skills/reviewer": "2.3.1"
checksums:
  "registry.internal/agent-skills/reviewer@2.3.1": "sha256:..."
activation_gates:
  require_human_approval_for:
    - "github.write"
    - "bash.exec"
```

Interpretation: - Discovery scope is constrained first. - Install/import are pinned and integrity-checked. - High-impact capabilities require explicit activation approval.

## 5.8 Key Takeaways

- Treat **artefact**, **discovery**, **import**, **install**, and **activate** as distinct terms.
- Discovering a tool endpoint is not the same as activating its capabilities.
- Use taxonomy-first language: tool, skill, agent, and workflow fragment are different artefact types.
- Prefer explicit, pinned configuration when provenance or authenticity is uncertain.
- Apply controls by stage: allowlist at discovery, verify at import/install, and least privilege at activation/runtime.

# Chapter 6

# GitHub Agentic Workflows (GH-AW)

## 6.1 Chapter Preview

- Explain how GH-AW compiles markdown into deterministic workflows.
- Show how to set up GH-AW with the supported setup actions.
- Highlight safety controls (permissions, safe outputs, approvals).

## 6.2 Why GH-AW Matters

GitHub Agentic Workflows (GH-AW) (<https://github.github.io/gh-aw/>) turns natural language into automated repository agents that run inside GitHub Actions. Instead of writing large YAML pipelines by hand, you write markdown instructions that an AI agent executes with guardrails. The result is a workflow you can read like documentation but run like automation.

At a glance, GH-AW provides:

- **Natural language workflows:** Markdown instructions drive the agent's behavior.
- **Compile-time structure:** Markdown is compiled into GitHub Actions workflows for reproducibility.
- **Security boundaries:** Permissions, tools, and safe outputs define what the agent can and cannot do.
- **Composable automation:** Imports and shared components enable reuse across repositories.

## 6.3 Core Workflow Structure

A GH-AW workflow is a markdown file with frontmatter and instructions:

```
---
on:
  issues:
    types: [opened]
permissions:
  contents: read
tools:
  edit:
    github:
      toolsets: [issues]
engine: copilot
---

# Triage this issue
Read issue $${{ github.event.issue.number }} and summarize it.
```

## Key parts:

1. **Frontmatter**
  - `on`: GitHub Actions triggers (issues, schedules, dispatch, etc.)
  - `permissions`: least-privilege access to GitHub APIs
  - `tools`: the capabilities your agent can invoke (edit, bash, web, github)
  - `engine`: AI model/provider (Copilot, Claude Code, Codex)
2. **Markdown instructions**
  - Natural language steps for the agent
  - Context variables from the event payload (issue number, PR, repo)

## 6.4 How GH-AW Runs

GH-AW compiles markdown workflows into `.lock.yml` GitHub Actions workflows. The compiled file is what GitHub actually executes, but the markdown remains the authoritative source. This gives you readable automation with predictable execution.

### 6.4.1 File Location

Both the source markdown files and the compiled `.lock.yml` files live in the `.github/workflows/` directory:

```
.github/workflows/  
|-- triage.md      # Source (human-editable)  
|-- triage.lock.yml # Compiled (auto-generated, do not edit)  
|-- docs-refresh.md  
`-- docs-refresh.lock.yml
```

Use `gh aw compile` (from the GH-AW CLI at <https://github.com/github/gh-aw>) in your repository root to generate `.lock.yml` files from your markdown sources. Only edit the `.md` files; the `.lock.yml` files are regenerated on compile.

If you do not vendor the GH-AW `actions/` directory in your repository, you can instead reference the upstream setup action directly (pin to a commit SHA for security):

```
- name: Setup GH-AW scripts  
  uses: github/gh-aw/actions/setup@5a4d651e3bd33de46b932d898c20c5619162332e  
  with:  
    destination: /opt/gh-aw/actions
```

### 6.4.2 Key Behaviors

- **Frontmatter edits require recompile.**
- **Markdown instruction updates can often be edited directly** (the runtime loads the markdown body).
- **Shared components** can be stored as markdown files without `on:`; they are imported, not compiled.

## 6.5 Compilation Model Examples

GH-AW compilation is mostly a structural translation: frontmatter becomes the workflow header, the markdown body is packaged as a script or prompt payload, and imports are inlined or referenced. The compiled `.lock.yml` is the contract GitHub Actions executes. The examples below show how a markdown workflow turns into a compiled job.

### 6.5.1 Example 1: Issue Triage Workflow

Source markdown (`.github/workflows/triage.md`)

```
---
on:
  issues:
    types: [opened]
permissions:
  contents: read
  issues: write
tools:
  github:
    toolsets: [issues]
engine: copilot
---

# Triage this issue
Read issue ${github.event.issue.number} and summarize it.
Then add labels: needs-triage and needs-owner.
```

Compiled workflow (`.github/workflows/triage.lock.yml`)

```
name: GH-AW triage
on:
  issues:
    types: [opened]
permissions:
  contents: read
  issues: write
jobs:
  agent:
    agent:
      runs-on: ubuntu-latest
    steps:
      - name: Checkout actions folder
        uses: actions/checkout@8e8c483db84b4bee98b60c0593521ed34d9990e8 # v6
        with:
          sparse-checkout: |
            actions
          persist-credentials: false
      - name: Setup GH-AW scripts
        uses: ./actions/setup
        with:
          destination: /opt/gh-aw/actions
      - name: Run GH-AW agent (generated)
        uses: actions/github-script@ed597411d8f924073f98dfc5c65a23a2325f34cd # v8.0.0
        with:
          script: |
            const { setupGlobals } = require('/opt/gh-aw/actions/setup_globals.cjs');
            setupGlobals(core, github, context, exec, io);
            // Generated execution script omitted for brevity.
```

#### What changed during compilation

- Frontmatter was converted into workflow metadata (`on`, `permissions`, `jobs`).
- Generated steps reference the GH-AW scripts copied by the setup action.
- The markdown body became the prompt payload executed by the agent runtime.

### 6.5.2 Example 2: Reusable Component + Import

Component (.github/workflows/shared/common-tools.md)

```
---
```

```
tools:
  bash:
  edit:
engine: copilot
---
```

Workflow using an import (.github/workflows/docs-refresh.md)

```
---
```

```
on:
  workflow_dispatch:
permissions:
  contents: write
imports:
  - shared/common-tools.md
---
```

```
# Refresh docs
Update the changelog with the latest release notes.
```

Compiled workflow (.github/workflows/docs-refresh.lock.yml)

```
name: GH-AW docs refresh
on:
  workflow_dispatch:
permissions:
  contents: write
jobs:
  agent:
    agent:
      runs-on: ubuntu-latest
    steps:
      - name: Checkout actions folder
        uses: actions/checkout@8e8c483db84b4bee98b60c0593521ed34d9990e8 # v6
        with:
          sparse-checkout: |
            actions
          persist-credentials: false
      - name: Setup GH-AW scripts
        uses: ./actions/setup
        with:
          destination: /opt/gh-aw/actions
      - name: Run GH-AW agent (generated)
        uses: actions/github-script@ed597411d8f924073f98dfc5c65a23a2325f34cd # v8.0.0
        with:
          script: |
            const { setupGlobals } = require('/opt/gh-aw/actions/setup_globals.cjs');
            setupGlobals(core, github, context, exec, io);
            // Generated execution script omitted for brevity.
```

What changed during compilation

- imports were resolved and merged with the workflow frontmatter.
- The component's tools and engine were applied to the final workflow.
- Only workflows with on: are compiled; components remain markdown-only.

### 6.5.3 Example 3: Safe Outputs in the Compiled Job

Source markdown (`.github/workflows/release-notes.md`)

```
---
on:
  workflow_dispatch:
permissions:
  contents: read
  pull-requests: write
tools:
  edit:
safe_outputs:
  pull_request_body:
    format: markdown
engine: copilot
---

# Draft release notes
Summarize commits since the last tag and open a PR with the notes.
```

Compiled workflow (`.github/workflows/release-notes.lock.yml`)

```
name: GH-AW release notes
on:
  workflow_dispatch:
permissions:
  contents: read
  pull-requests: write
jobs:
  agent:
    agent:
      runs-on: ubuntu-latest
    steps:
      - name: Checkout actions folder
        uses: actions/checkout@8e8c483db84b4bee98b60c0593521ed34d9990e8 # v6
        with:
          sparse-checkout: |
            actions
          persist-credentials: false
      - name: Setup GH-AW scripts
        uses: ./actions/setup
        with:
          destination: /opt/gh-aw/actions
      - name: Run GH-AW agent (generated)
        uses: actions/github-script@ed597411d8f924073f98dfc5c65a23a2325f34cd # v8.0.0
        with:
          script: |
            const { setupGlobals } = require('/opt/gh-aw/actions/setup_globals.cjs');
            setupGlobals(core, github, context, exec, io);
            // Generated execution script omitted for brevity.
```

#### What changed during compilation

- `safe_outputs` was translated into the generated safe-output scripts invoked by the job.
- The prompt stayed identical; guardrails are enforced by the compiled job.

## 6.6 Tools, Safe Inputs, and Safe Outputs

GH-AW workflows are designed for safety by default. Agents run with minimal access and must declare tools explicitly.

**Warning:** Treat CI secrets and tokens as production credentials. Use least-privilege permissions, require human approval for write actions, and keep all agent actions auditable.

### 6.6.1 Tools

Tools are capabilities the agent can use:

- **edit**: modify files in the workspace
- **bash**: run shell commands (by default only safe commands)
- **web-fetch / web-search**: fetch or search web content
- **github**: operate on issues, PRs, discussions, projects
- **playwright**: browser automation for UI checks

### 6.6.2 Safe Outputs

Write actions (creating issues, comments, commits) can be routed through safe outputs to sanitize what the agent writes. This keeps the core job read-only and limits accidental changes.

### 6.6.3 Safe Inputs

You can define safe inputs to structure what the agent receives. This is a good place to validate schema-like data for tools or commands.

## 6.7 Imports and Reusable Components

For terminology and trust-model definitions, see Discovery and Imports. This section focuses only on GH-AW-specific syntax and composition patterns.

GH-AW supports imports in two ways:

- **Frontmatter imports**

```
imports:  
  - shared/common-tools.md  
  - shared/research-library.md
```

- **Markdown directive**

```
{% raw %}{{#import shared/common-tools.md}}{% endraw %}
```

In GH-AW, these imports are typically workflow-fragment artefacts: shared prompts, tool declarations, and policy snippets. Keep reusable fragments in files without `on:` so they can be imported as components rather than compiled as standalone workflows.

## 6.8 ResearchPlanAssign: A Pattern for Self-Maintaining Books

GH-AW documents a **ResearchPlanAssign** strategy: a scaffolded loop that keeps humans in control while delegating research and execution to agents.

**Phase 1: Research** - A scheduled agent scans the repo or ecosystem (new libraries, frameworks, scaffolds).  
- It produces a report in an issue or discussion.

**Phase 2: Plan** - Maintainers review the report and decide whether to proceed. - A planning agent drafts the implementation steps if approved.

**Phase 3: Assign & Implement** - Agents are assigned to implement the approved changes. - Updates are validated, committed, and published.

This pattern maps well to this book: use scheduled research to discover new agentic tooling, post a proposal issue, build consensus, then update the chapters and blog.

## 6.9 Applying GH-AW to This Repository

Here's how GH-AW can drive the book's maintenance:

1. **Research (scheduled)**
  - Use web-search tooling to scan for new agentic workflow libraries.
  - Produce a structured report in a GitHub issue.
2. **Consensus (issues/discussions)**
  - Collect votes or comments to accept/reject the proposal.
  - Label outcomes (`accepted`, `needs-revision`, `rejected`).
3. **Implementation (assigned agent)**
  - Update or add chapters.
  - Refresh the table of contents and homepage.
  - Add a blog post summarizing the update.
4. **Publish (automation)**
  - Pages and PDF rebuild automatically after merge.

This approach keeps the book aligned with the latest GH-AW practices while maintaining a transparent, auditable workflow.

## 6.10 Key Takeaways

- GH-AW turns markdown instructions into reproducible GitHub Actions workflows.
- Frontmatter defines triggers, permissions, tools, and models.
- Imports enable composable, reusable workflow building blocks.
- Safe inputs/outputs and least-privilege permissions reduce risk.
- ResearchPlanAssign provides a practical loop for continuous, agent-powered improvement.

# Chapter 7

## GitHub Agents

### 7.1 Chapter Preview

- Describe how agents operate inside GitHub issues, PRs, and Actions.
- Show safe assignment, review, and approval flows.
- Map GitHub agent capabilities to real repository workflows.

### 7.2 Understanding GitHub Agents

GitHub Agents represent a new paradigm in software development automation. They are AI-powered assistants that can understand context, make decisions, and take actions within the GitHub ecosystem. Unlike traditional automation that follows predefined scripts, agents can adapt to situations, reason about problems, and collaborate with humans and other agents.

This chapter explores the landscape of GitHub Agents, their capabilities, and how to leverage them effectively in your development workflows.

### 7.3 The GitHub Agent Ecosystem

#### 7.3.1 GitHub Copilot

GitHub Copilot (<https://docs.github.com/en/copilot>) is the foundation of GitHub's AI-powered development tools. It provides:

- **Code Completion:** Real-time suggestions as you type
- **Chat Interface:** Natural language conversations about code
- **Context Awareness:** Understanding of your codebase and intent

```
# Example: Copilot helping write a function
# Just start typing a comment describing what you need:
# Function to validate email addresses using regex
def validate_email(email):
    import re
    pattern = r'^[a-zA-Z0-9._%+-]+@[a-zA-Z0-9.-]+\.[a-zA-Z]{2,}$'
    return bool(re.match(pattern, email))
```

### 7.3.2 GitHub Copilot Coding Agent

The Coding Agent extends Copilot's capabilities to autonomous task completion. See <https://docs.github.com/en/copilot/how-tos/use-copilot-agents/coding-agent> for the supported assignment and review flow:

- **Assigned Tasks:** Receives issues or requests and works independently
- **Multi-File Changes:** Can modify multiple files across a codebase
- **Pull Request Creation:** Generates complete PRs with descriptions
- **Iterative Development:** Responds to review feedback

#### Key Characteristics:

| Feature  | Description                                  |
|----------|--|
| Autonomy | Works independently on assigned tasks        |
| Scope    | Can make changes across entire repositories  |
| Output   | Creates branches, commits, and pull requests |
| Review   | All changes go through normal PR review      |

### 7.3.3 GitHub Actions Agents

Agents can be orchestrated through GitHub Actions workflows:

```
name: Agent Workflow
on:
  issues:
    types: [opened]

jobs:
  agent-task:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v6
      - name: Process with Agent
        uses: actions/github-script@v8
        with:
          script: |
            // Agent logic to analyze and respond
            const issue = context.payload.issue;
            // ... agent processing
```

**Tip:** In production workflows, pin third-party actions to a full commit SHA to reduce supply-chain risk.

**Warning:** Require human approval before any agent-created PR is merged, and log all agent actions for auditability.

## 7.4 Agent Capabilities

### 7.4.1 Reading and Understanding

Agents can read and understand:

- **Code:** Source files, configurations, dependencies
- **Documentation:** READMEs, wikis, comments
- **Issues and PRs:** Descriptions, comments, reviews
- **Repository Structure:** File organization, patterns

## 7.4.2 Writing and Creating

Agents can produce:

- **Code Changes:** New files, modifications, refactoring
- **Documentation:** READMEs, comments, API docs
- **Issues and Comments:** Status updates, analysis reports
- **Pull Requests:** Complete PRs with proper descriptions

## 7.4.3 Reasoning and Deciding

Agents can:

- **Analyze Problems:** Understand issue context and requirements
- **Plan Solutions:** Break down tasks into steps
- **Make Decisions:** Choose between approaches
- **Adapt:** Respond to feedback and changing requirements

# 7.5 Multi-Agent Orchestration

## 7.5.1 Why Multiple Agents?

Single agents have limitations. Multi-agent systems provide:

1. **Specialization:** Each agent excels at specific tasks
2. **Perspective Diversity:** Different models bring different strengths
3. **Scalability:** Parallel processing of independent tasks
4. **Resilience:** Failure of one agent doesn't stop the workflow

## 7.5.2 Orchestration Patterns

### 7.5.2.1 Sequential Pipeline

Agents work in sequence, each building on the previous:

```
Issue -> ACK Agent -> Research Agent -> Writer Agent -> Review Agent -> Complete
```

Example workflow stages:

```
jobs:  
  stage-1-acknowledge:  
    runs-on: ubuntu-latest  
    if: github.event.action == 'opened'  
    # Acknowledge and validate  
  
  stage-2-research:  
    runs-on: ubuntu-latest  
    needs: stage-1-acknowledge  
    if: needs.stage-1-acknowledge.outputs.is_relevant == 'true'  
    # Research and analyze  
  
  stage-3-write:  
    runs-on: ubuntu-latest  
    needs: stage-2-research  
    # Create content
```

### 7.5.2.2 Parallel Discussion

Multiple agents contribute perspectives simultaneously:

```
jobs:
  discuss:
    runs-on: ubuntu-latest
    strategy:
      matrix:
        agent: [claude, copilot, gemini]
    steps:
      - name: Agent Perspective
        # Each agent provides its view
```

### 7.5.2.3 Human-in-the-Loop

Agents work until human decision is needed:

```
Agents work -> Human checkpoint -> Agents continue
```

This pattern is essential for: - Approving significant changes - Resolving ambiguous decisions - Quality assurance

### 7.5.3 Agent Handoff Protocol

When agents need to pass context to each other:

1. **State in Labels:** Use GitHub labels to track workflow stage
2. **Context in Comments:** Agents document their findings in issue comments
3. **Structured Output:** Use consistent formats for machine readability

```
# Example: Structured agent output
- name: Agent Report
  uses: actions/github-script@v8
  with:
    script: |
      const report = {
        stage: 'research',
        findings: [...],
        recommendation: 'proceed',
        nextAgent: 'writer'
      };
      // Store in comment or labels
```

## 7.6 Implementing GitHub Agents

### 7.6.1 Agent Definition Files

Define agents in markdown files with frontmatter:

```
---
name: Research Agent
description: Analyzes issues and researches documentation
tools:
  github:
    toolsets: [issues]
  web-search:
  edit:
```

```

---
# Research Agent

You are the research agent for this repository.
Your role is to analyze suggestions and assess their value.

## Tasks
1. Search existing documentation
2. Find relevant external sources
3. Assess novelty and interest
4. Report findings

```

### 7.6.2 Agent Configuration

Control agent behavior through configuration:

```

# .github/agents/config.yml
agents:
  research:
    enabled: true
    model: copilot
    timeout: 300

  writer:
    enabled: true
    model: copilot
    requires_approval: true

  safety:
    max_file_changes: 10
    protected_paths:
      - .github/workflows/
      - SECURITY.md

```

### 7.6.3 Error Handling

Agents should handle failures gracefully:

```

- name: Agent Task with Error Handling
  id: agent_task
  continue-on-error: true
  uses: actions/github-script@v8
  with:
    script: |
      try {
        // Agent logic
      } catch (error) {
        await github.rest.issues.addLabels({
          owner: context.repo.owner,
          repo: context.repo.repo,
          issue_number: context.issue.number,
          labels: ['agent-error']
        });
        await github.rest.issues.createComment({
          owner: context.repo.owner,
          repo: context.repo.repo,
          issue_number: context.issue.number,
        })
      }

```

```
        body: `WARNING: Agent encountered an error: ${error.message}`  
    );  
}
```

## 7.7 Best Practices

### 7.7.1 Clear Agent Personas

Give each agent a clear identity and responsibility:

```
## You Are: The Research Agent  
  
**Your Role:** Investigate and analyze  
**You Are Not:** A decision maker or implementer  
**Hand Off To:** Writer Agent after research is complete
```

### 7.7.2 Structured Communication

Use consistent formats for agent-to-agent communication:

```
## Agent Report Format  
  
### Status: [Complete/In Progress/Blocked]  
### Findings:  
- Finding 1  
- Finding 2  
### Recommendation: [Proceed/Revise/Decline]  
### Next Stage: [stage-name]
```

### 7.7.3 Human Checkpoints

Always include human review points:

- Before significant changes
- After agent recommendations
- Before closing issues

### 7.7.4 Audit Trail

Maintain visibility into agent actions:

- All agent actions should be visible in comments
- Use labels to track workflow state
- Log important decisions and reasoning

### 7.7.5 Graceful Degradation

Design for agent failures:

- Use `continue-on-error` for non-critical steps
- Provide manual fallback options
- Alert maintainers when intervention is needed

## 7.8 Security Considerations

### 7.8.1 Least Privilege

Agents should have minimal permissions:

```
permissions:  
  contents: read  # Only write if needed  
  issues: write   # To comment and label  
  pull-requests: write # Only if creating PRs
```

### 7.8.2 Input Validation

Validate data before agent processing:

```
// Validate issue body before processing  
const body = context.payload.issue.body || '';  
if (body.length > 10000) {  
  throw new Error('Issue body too long');  
}
```

### 7.8.3 Output Sanitization

Sanitize agent outputs:

```
// Escape user content in agent responses  
const safeTitle = issueTitle.replace(/<[^>]>/g, ''');
```

### 7.8.4 Protected Resources

Prevent agents from modifying sensitive files:

```
# In workflow: check protected paths  
- name: Check Protected Paths  
  run: |  
    CHANGED_FILES=$(git diff --name-only HEAD~1)  
    if echo "$CHANGED_FILES" | grep -E "^(SECURITY|\.github/workflows/)" ; then  
      echo "Protected files modified - requires human review"  
      exit 1  
    fi
```

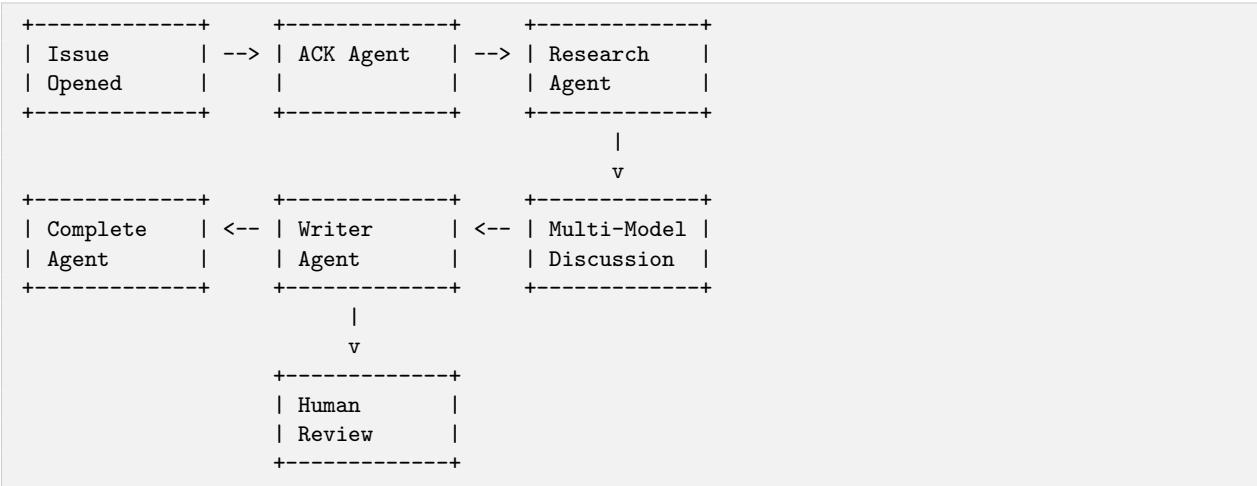
## 7.9 Real-World Example: This Book

This very book uses GitHub Agents for self-maintenance:

### 7.9.1 The Multi-Agent Workflow

1. **ACK Agent:** Acknowledges new issue suggestions
2. **Research Agent:** Analyzes novelty and relevance
3. **Claude Agent:** Provides safety and clarity perspective
4. **Copilot Agent:** Provides developer experience perspective
5. **Writer Agent:** Drafts new content
6. **Completion Agent:** Finalizes and closes issues

## 7.9.2 How It Works



## 7.9.3 Configuration

The workflow is defined using GitHub Agentic Workflows (GH-AW). The repository includes:

- GH-AW workflows: `.github/workflows/issue-*.lock.yml`
- Agent definitions: `.github/agents/*.md`

For a detailed explanation of the workflow architecture and why GH-AW is the canonical approach, see the repository's `WORKFLOWS.md` documentation.

## 7.10 Multi-Agent Platform Compatibility

Modern repositories need to support multiple AI agent platforms. Different coding assistants—GitHub Copilot (<https://docs.github.com/en/copilot>), Claude (<https://code.claude.com/docs>), OpenAI Codex (<https://openai.com/index/introducing-codex/>), and others—each have their own ways of receiving project-specific instructions. This section explains how to structure a repository for cross-platform agent compatibility.

### 7.10.1 The Challenge of Agent Diversity

When multiple AI agents work with your repository, you face a coordination challenge:

- **GitHub Copilot** reads `.github/copilot-instructions.md` for project-specific guidance
- **Claude** automatically incorporates `CLAUDE.md`; `AGENTS.md` may still be useful as shared project documentation, but should be explicitly referenced for reliable Claude workflows
- **OpenAI Codex (GPT-5.2-Codex)** can be configured with system instructions and skills packaged via `SKILL.md` (see <https://platform.openai.com/docs/guides/codex/skills>)
- **Generic agents** look for `AGENTS.md` as the emerging standard

Each platform has slightly different expectations, but the core information they need is similar.

### 7.10.2 Repository Documentation as Agent Configuration

Your repository's documentation files serve dual purposes—they guide human contributors AND configure AI agents. Key files include:

| File                         | Human Purpose           | Agent Purpose                          |
|------------------------------|-------------------------|--|
| <code>README.md</code>       | Project overview        | Context for understanding the codebase |
| <code>CONTRIBUTING.md</code> | Contribution guidelines | Workflow rules and constraints         |

| File                            | Human Purpose | Agent Purpose                  |
|---------------------------------|---------------|--------------------------------|
| .github/copilot-instructions.md | N/A           | Copilot-specific configuration |
| AGENTS.md                       | N/A           | Generic agent instructions     |
| CLAUDE.md                       | N/A           | Claude-specific configuration  |

### 7.10.3 The copilot-instructions.md File

GitHub Copilot reads `.github/copilot-instructions.md` to understand how to work with your repository. This file should include:

```
# Copilot Instructions for [Project Name]

## Project Overview
Brief description of what this project does.

## Tech Stack
- **Language**: Python 3.11
- **Framework**: FastAPI
- **Database**: PostgreSQL
- **Testing**: pytest

## Coding Guidelines
- Follow PEP 8 style guide
- All functions require type hints
- Tests are required for new features

## File Structure
Describe important directories and their purposes.

## Build and Test Commands
- `make test` - Run all tests
- `make lint` - Run linters
- `make build` - Build the project

## Important Constraints
- Never commit secrets or credentials
- Protected paths: .github/workflows/, secrets/
- All PRs require review before merge
```

### 7.10.4 Cross-Platform Strategy

For maximum compatibility across AI agent platforms, follow these practices:

1. **Pick a canonical source per platform** (`AGENTS.md` for many coding agents, `CLAUDE.md` for Claude)
2. **Cross-reference shared guidance** between platform files to reduce drift
3. **Keep instructions DRY** by avoiding unnecessary duplication
4. **Test with multiple agents** to ensure instructions work correctly

Example hierarchy:

```
project/
|-- AGENTS.md                      # Canonical agent instructions
|-- CLAUDE.md                        # Claude-specific (may reference AGENTS.md)
|-- .github/
|   '-- copilot-instructions.md     # Copilot-specific (may reference AGENTS.md)
```

```
-- src/
  -- AGENTS.md          # Module-specific instructions
```

### 7.10.5 This Repository's Approach

This book repository demonstrates multi-platform compatibility:

- `.github/copilot-instructions.md` - Detailed Copilot configuration with project structure, coding guidelines, and constraints
- **Skills and Tools Management** and **Agents for Coding** discuss AGENTS.md as the emerging standard
- **Documentation files** (README, CONTRIBUTING, etc.) provide context any agent can use
- **GH-AW workflows** use the `engine: copilot` setting but the pattern works with other engines

The key insight is that well-structured documentation benefits both human developers and AI agents. When you write clear README files, contribution guidelines, and coding standards, you're simultaneously creating better agent configuration.

### 7.10.6 Best Practices for Agent-Friendly Repositories

1. **Be explicit about constraints:** Clearly state what agents should NOT do
2. **Document your tech stack:** Agents perform better when they understand the tools in use
3. **Describe the project structure:** Help agents navigate your codebase efficiently
4. **Provide examples:** Show preferred patterns through code examples
5. **List protected paths:** Specify files agents should not modify
6. **Include build/test commands:** Enable agents to verify their changes
7. **State coding conventions:** Help agents write consistent code

## 7.11 Future of GitHub Agents

### 7.11.1 Emerging Capabilities

- **Code Generation:** Agents writing production-quality code
- **Test Authoring:** Automatic test creation and maintenance
- **Documentation Sync:** Keeping docs in sync with code
- **Security Analysis:** Proactive vulnerability detection

### 7.11.2 Integration Trends

- **IDE Integration:** Deeper VS Code and editor integration
- **CI/CD Native:** Agents as first-class CI/CD citizens
- **Cross-Repo:** Agents working across multiple repositories
- **Multi-Cloud:** Agents coordinating across platforms

## 7.12 Key Takeaways

1. **GitHub Agents** are AI-powered assistants that can reason, decide, and act within repositories.
2. **Copilot Coding Agent** can autonomously complete tasks and create pull requests.
3. **Multi-agent orchestration** enables specialized, resilient, and scalable automation.
4. **Human checkpoints** remain essential for quality and safety.
5. **Clear protocols** for agent communication ensure smooth handoffs.

6. **Security** must be designed into agent workflows from the start.
7. **Multi-platform compatibility** is achieved through well-structured documentation (copilot-instructions.md, AGENTS.md, etc.).
8. **This book** demonstrates these concepts through its own multi-agent maintenance workflow.

## 7.13 Learn More

### 7.13.1 Repository Documentation

This book's repository includes comprehensive documentation that demonstrates OSS best practices:

- **README** - Overview and quick start guide
- **CONTRIBUTING** - How to contribute using the multi-agent workflow
- **WORKFLOWS** - Detailed workflow guide and GH-AW explanation
- **SETUP** - Installation and configuration instructions
- **WORKFLOW\_PLAYBOOK** - Agentic workflow maintenance patterns
- **PROJECT\_SUMMARY** - Complete project overview
- **SECURITY\_SUMMARY** - Security practices and scan results
- **CHANGELOG** - Version history and changes
- **CODE\_OF\_CONDUCT** - Community guidelines
- **LICENSE** - MIT License

### 7.13.2 Agent Configuration Files

These files configure how AI agents work with this repository:

- **.github/copilot-instructions.md** - GitHub Copilot-specific configuration including project structure, coding guidelines, and constraints

These documents serve as both useful references and examples of how to structure documentation for projects using agentic workflows.

### 7.13.3 Related Sections

- **Skills and Tools Management** - Covers AGENTS.md standard and MCP protocol for tool management
  - **GitHub Agentic Workflows (GH-AW)** - GH-AW specification and engine configuration
  - **Agents for Coding** - Detailed coverage of coding agent platforms
-

# Chapter 8

## Agents for Coding

### 8.1 Chapter Preview

- Compare coding-agent architectures and team patterns.
- Show the correct way to configure GitHub Copilot coding agent.
- Provide buildable examples and clear labels for pseudocode.

### 8.2 Introduction

Coding agents represent the most mature category of AI agents in software development. They have evolved from simple autocomplete tools to autonomous entities capable of planning, writing, testing, debugging, and even scaffolding entire software architectures with minimal human input. This chapter explores the specialized architectures, scaffolding patterns, and best practices for deploying agents in coding workflows.

### 8.3 The Evolution of Coding Agents

#### 8.3.1 From Autocomplete to Autonomy

The progression of coding agents follows a clear trajectory:

1. **Code Completion (2020-2022)**: Basic pattern matching and next-token prediction
2. **Context-Aware Assistance (2022-2024)**: Understanding project structure and intent
3. **Task-Oriented Agents (2024-present)**: Completing multi-step tasks independently
4. **Autonomous Development (emerging)**: Full feature implementation, testing, and deployment

#### 8.3.2 Current Capabilities

Modern coding agents can:

- **Understand Requirements**: Parse natural language specifications and translate them to code
- **Plan Solutions**: Break down complex features into implementable steps
- **Generate Code**: Write production-quality code across multiple files
- **Test and Debug**: Create tests, identify bugs, and fix issues
- **Scaffold Projects**: Initialize projects with appropriate structure and configuration
- **Review and Refactor**: Analyze code quality and suggest improvements

## 8.4 Specialized Architectures

### 8.4.1 Single-Agent Architectures

The simplest architecture involves one agent with access to all necessary tools.

Example 7-2. Single-agent architecture (pseudocode)

```
class CodingAgent:
    """Single-agent architecture for coding tasks"""

    def __init__(self, llm, tools):
        self.llm = llm
        self.tools = {
            'file_read': FileReadTool(),
            'file_write': FileWriteTool(),
            'terminal': TerminalTool(),
            'search': CodeSearchTool(),
            'test_runner': TestRunnerTool()
        }
        self.context = AgentContext()

    async def execute(self, task: str) -> dict:
        """Execute a coding task end-to-end"""
        # 1. Understand the task
        plan = await self.plan_task(task)

        # 2. Execute each step
        results = []
        for step in plan.steps:
            result = await self.execute_step(step)
            results.append(result)

        # Adapt based on results
        if not result.success:
            plan = await self.replan(plan, result)

        return {'success': True, 'results': results}
```

**Best for:** Simple tasks, small codebases, single-developer workflows.

### 8.4.2 Multi-Agent Architectures

Complex projects benefit from specialized agents working together.

Example 7-3. Multi-agent architecture (pseudocode)

```
class CodingAgentTeam:
    """Multi-agent architecture mirroring a development team"""

    def __init__(self):
        self.architect = ArchitectAgent()
        self.implementer = ImplementerAgent()
        self.tester = TesterAgent()
        self.reviewer = ReviewerAgent()
        self.coordinator = CoordinatorAgent()

    async def execute_feature(self, specification: str):
        """Execute a feature request using the agent team"""

```

```

# 1. Architecture phase
design = await self.architect.design(specification)

# 2. Implementation phase (can be parallelized)
implementations = await asyncio.gather(*[
    self.implementer.implement(component)
    for component in design.components
])

# 3. Testing phase
test_results = await self.tester.test(implementations)

# 4. Review phase
review = await self.reviewer.review(implementations)

# 5. Iteration if needed
if not review.approved:
    return await self.handle_review_feedback(review)

return {'success': True, 'implementation': implementations}

```

**Best for:** Large projects, team environments, complex features.

#### 8.4.3 Subagent and Swarms Mode

Modern frameworks like Claude Code support dynamic subagent spawning:

```

class SwarmCoordinator:
    """Coordinate a swarm of specialized subagents"""

    def __init__(self, max_agents=10):
        self.max_agents = max_agents
        self.active_agents = {}

    async def spawn_subagent(self, task_type: str, context: dict):
        """Spawn a specialized subagent for a specific task"""

        agent_configs = {
            'frontend': FrontendAgentConfig(),
            'backend': BackendAgentConfig(),
            'devops': DevOpsAgentConfig(),
            'security': SecurityAgentConfig(),
            'documentation': DocsAgentConfig()
        }

        config = agent_configs.get(task_type)
        agent = await self.create_agent(config, context)

        self.active_agents[agent.id] = agent
        return agent

    async def execute_parallel(self, tasks: list):
        """Execute multiple tasks in parallel using subagents"""

        agents = [
            await self.spawn_subagent(task.type, task.context)
            for task in tasks
        ]

```

```

        results = await asyncio.gather(*[
            agent.execute(task)
            for agent, task in zip(agents, tasks)
        ])

        return self.aggregate_results(results)
    
```

## 8.5 Scaffolding for Coding Agents

### 8.5.1 Project Initialization

Coding agents need scaffolding that helps them understand and work with projects:

```

# .github/agents/coding-agent.yml
name: coding-agent
description: Scaffolding for coding agent operations

workspace:
  root: ../
  source_dirs: [src/, lib/]
  test_dirs: [tests/, spec/]
  config_files: [package.json, tsconfig.json, .eslintrc]

conventions:
  language: typescript
  framework: express
  testing: jest
  style: prettier + eslint

tools:
  enabled:
    - file_operations
    - terminal
    - git
    - package_manager
  restricted:
    - network_access
    - system_commands

safety:
  max_file_changes: 20
  protected_paths:
    - .github/workflows/
    - .env*
    - secrets/
  require_tests: true
  require_review: true

```

### 8.5.2 The AGENTS.md Standard

The **AGENTS.md** file has emerged as the de facto standard for providing AI coding agents with project-specific instructions. It serves as a “README for agents,” offering structured, machine-readable guidance.

#### 8.5.2.1 Purpose and Placement

```
project/
```

```

|-- AGENTS.md          # Root-level agent instructions
|-- src/
|   |-- AGENTS.md      # Module-specific instructions
|-- tests/
|   |-- AGENTS.md      # Testing conventions
`-- docs/
    |-- AGENTS.md      # Documentation guidelines

```

Agents use the nearest AGENTS.md file, enabling scoped configuration for monorepos or complex projects.

#### 8.5.2.2 Standard Structure

```

# AGENTS.md

## Project Overview
Brief description of the project and its purpose.

## Setup Instructions
How to install dependencies and configure the environment.

## Coding Conventions
- Language: TypeScript 5.x
- Style guide: Airbnb
- Formatting: Prettier with provided config

## Testing Requirements
- Framework: Jest
- Coverage threshold: 80%
- Required test types: unit, integration

## Build and Deploy
- Build command: `npm run build`
- Deploy process: CI/CD via GitHub Actions

## Agent-Specific Notes
- Always run linting before committing
- Never modify files in `vendor/`
- Secrets are in environment variables, never hardcoded

```

#### 8.5.2.3 Benefits

1. **Consistency:** All agents receive the same instructions
2. **Onboarding:** New agents (or new agent sessions) understand the project immediately
3. **Safety:** Clear boundaries prevent accidental damage
4. **Maintainability:** Single source of truth for agent behavior

#### 8.5.3 Context Management

Coding agents need effective context management to work across large codebases:

```

class CodingContext:
    """Manage context for coding agents"""

    def __init__(self, workspace_root: str):
        self.workspace_root = workspace_root
        self.file_index = FileIndex(workspace_root)
        self.symbol_table = SymbolTable()

```

```

        self.active_files = LRUCache(max_size=50)

    def get_relevant_context(self, task: str) -> dict:
        """Get context relevant to the current task"""

        # 1. Parse task to identify relevant files/symbols
        entities = self.extract_entities(task)

        # 2. Retrieve relevant files
        files = self.file_index.search(entities)

        # 3. Get symbol definitions
        symbols = self.symbol_table.lookup(entities)

        # 4. Include recent changes
        recent = self.get_recent_changes()

        return {
            'files': files,
            'symbols': symbols,
            'recent_changes': recent,
            'workspace_config': self.get_config()
        }

    def update_context(self, changes: list):
        """Update context after agent makes changes"""
        for change in changes:
            self.file_index.update(change.path)
            self.symbol_table.reindex(change.path)
            self.active_files.add(change.path)

```

#### 8.5.4 Tool Registries

Coding agents need well-organized tool access:

```

class CodingToolRegistry:
    """Registry of tools available to coding agents"""

    def __init__(self):
        self._tools = {}
        self._register_default_tools()

    def _register_default_tools(self):
        """Register standard coding tools"""

        # File operations
        self.register('read_file', ReadFileTool())
        self.register('write_file', WriteFileTool())
        self.register('search_files', SearchFilesTool())

        # Code operations
        self.register('parse_ast', ParseASTTool())
        self.register('refactor', RefactorTool())
        self.register('format_code', FormatCodeTool())

        # Testing
        self.register('run_tests', RunTestsTool())
        self.register('coverage', CoverageTool())

```

```

# Git operations
self.register('git_status', GitStatusTool())
self.register('git_diff', GitDiffTool())
self.register('git_commit', GitCommitTool())

# Package management
self.register('npm_install', NpmInstallTool())
self.register('pip_install', PipInstallTool())

def get_tools_for_task(self, task_type: str) -> list:
    """Get tools appropriate for a task type"""

    task_tool_map = {
        'implementation': ['read_file', 'write_file', 'search_files', 'format_code'],
        'testing': ['read_file', 'write_file', 'run_tests', 'coverage'],
        'debugging': ['read_file', 'parse_ast', 'run_tests', 'git_diff'],
        'refactoring': ['read_file', 'write_file', 'parse_ast', 'refactor', 'run_tests']
    }

    tool_names = task_tool_map.get(task_type, list(self._tools.keys()))
    return [self._tools[name] for name in tool_names if name in self._tools]

```

## 8.6 Leading Coding Agent Platforms

### 8.6.1 GitHub Copilot and Coding Agent

GitHub Copilot has evolved from an IDE autocomplete tool to a full coding agent:

- **Copilot Chat:** Natural language interaction about code
- **Copilot Coding Agent:** Autonomous task completion and PR creation
- **Copilot Workspace:** Full development environment with agent integration

Copilot coding agent is not invoked via a custom `uses:` action. Instead, you assign work through GitHub Issues, Pull Requests, the agents panel, or by mentioning `@copilot`, and you customize its environment with a dedicated workflow file. See the official docs at <https://docs.github.com/en/copilot/how-tos/use-copilot-agents/coding-agent> and the environment setup guide at <https://docs.github.com/en/copilot/how-tos/use-copilot-agents/coding-agent/customize-the-agent-environment>.

Example 7-1. `.github/workflows/copilot-setup-steps.yml`

```

name: Copilot setup steps

on:
  push:
    paths:
      - .github/workflows/copilot-setup-steps.yml

jobs:
  # The job MUST be named copilot-setup-steps to be picked up by Copilot.
  copilot-setup-steps:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v6
      - name: Install dependencies
        run: |
          npm ci

```

## 8.6.2 Claude Code

Claude Code (<https://code.claude.com/docs>) provides multi-agent orchestration for complex development tasks:

- **Subagent Architecture:** Spawn specialized agents for different concerns
- **Swarms Mode:** Parallel execution of independent tasks
- **Extended Context:** Handle large codebases through intelligent context management

## 8.6.3 Cursor AI

Cursor (<https://www.cursor.com/>) is an AI-first code editor designed around agent workflows:

- **Project-Wide Understanding:** Indexes entire codebase for context
- **Multi-File Generation:** Creates and modifies multiple files in one operation
- **Framework Integration:** Deep understanding of popular frameworks

## 8.6.4 CodeGPT and Agent Marketplaces

CodeGPT (<https://codegpt.co/>) and marketplace-based approaches offer specialized agents:

- **Specialized Agents:** Over 200 pre-built agents for specific tasks
- **Custom Agent Creation:** Build and share domain-specific agents
- **Multi-Model Support:** Combine different LLMs for different tasks

## 8.7 Best Practices

### 8.7.1 Clear Task Boundaries

Define clear boundaries for what agents can and cannot do:

```
class TaskBoundary:  
    """Define boundaries for agent tasks"""  
  
    def __init__(self):  
        self.max_files = 20  
        self.max_lines_per_file = 500  
        self.timeout_seconds = 600  
        self.protected_patterns = [  
            r'\.env.*',  
            r'secrets/*',  
            r'\.github/workflows/*'  
        ]  
  
    def validate_task(self, task: dict) -> bool:  
        """Validate that a task is within boundaries"""  
        if len(task.get('files', [])) > self.max_files:  
            return False  
  
        for file_path in task.get('files', []):  
            if any(re.match(p, file_path) for p in self.protected_patterns):  
                return False  
  
        return True
```

### 8.7.2 Incremental Changes

Prefer small, focused changes over large rewrites:

```

class IncrementalChangeStrategy:
    """Strategy for making incremental changes"""

    def execute(self, large_change: Change) -> list:
        """Break large change into incremental steps"""

        # 1. Analyze the change
        components = self.decompose(large_change)

        # 2. Order by dependency
        ordered = self.topological_sort(components)

        # 3. Execute incrementally with validation
        results = []
        for component in ordered:
            result = self.apply_change(component)

            # Validate after each step
            if not self.validate(result):
                self.rollback(results)
                raise ChangeValidationError(result)

            results.append(result)

        return results

```

### 8.7.3 Test-Driven Development

Integrate testing into agent workflows:

```

class TDDAgent:
    """Agent that follows test-driven development"""

    async def implement_feature(self, specification: str):
        """Implement feature using TDD approach"""

        # 1. Write tests first
        tests = await self.generate_tests(specification)
        await self.write_tests(tests)

        # 2. Verify tests fail
        initial_results = await self.run_tests()
        assert not initial_results.all_passed

        # 3. Implement to pass tests
        implementation = await self.implement(specification, tests)

        # 4. Verify tests pass
        final_results = await self.run_tests()

        # 5. Refactor if needed
        if final_results.all_passed:
            await self.refactor_for_quality()

    return implementation

```

#### 8.7.4 Human Review Integration

Always include human checkpoints for significant changes:

Example 7-4. Human review workflow (pseudocode)

```
# Workflow with human review
name: Agent Implementation with Review
on:
  issues:
    types: [labeled]

jobs:
  implement:
    if: contains(github.event.issue.labels.*.name, 'agent-task')
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v6 # pin to a full SHA in production

      - name: Agent Implementation
        id: implement
        uses: ./actions/coding-agent

      - name: Create PR for Review
        uses: peter-evans/create-pull-request@v8 # pin to a full SHA in production
        with:
          title: "Agent: ${{ github.event.issue.title }}"
          body: |
            ## Agent Implementation

            This PR was created by an AI agent based on issue #${{ github.event.issue.number }}.

            **Please review carefully before merging.**
        labels: needs-human-review
        draft: true
```

**Note:** The `./actions/coding-agent` step is a placeholder for your organization's internal agent runner. Replace it with your approved agent execution mechanism.

## 8.8 Common Challenges

### 8.8.1 Context Window Limitations

Large codebases exceed agent context windows:

**Solution:** Implement intelligent context retrieval and summarization.

```
class ContextCompressor:
    """Compress context to fit within token limits"""

    def compress(self, files: list, max_tokens: int) -> str:
        """Compress file contents to fit token limit"""

        # Prioritize by relevance
        ranked = self.rank_by_relevance(files)

        # Include summaries for less relevant files
        context = []
        tokens_used = 0
```

```

for file in ranked:
    if tokens_used + file.tokens <= max_tokens:
        context.append(file.content)
        tokens_used += file.tokens
    else:
        # Include summary instead
        summary = self.summarize(file)
        context.append(summary)
        tokens_used += len(summary.split())

return '\n'.join(context)

```

### 8.8.2 Hallucination and Accuracy

Agents may generate plausible but incorrect code:

**Solution:** Implement validation and testing at every step.

### 8.8.3 Security Concerns

Agents with code access pose security risks:

**Solution:** Use sandboxing, permission scoping, and audit logging.

```

class SecureCodingEnvironment:
    """Secure environment for coding agent execution"""

    def __init__(self):
        self.sandbox = DockerSandbox()
        self.audit_log = AuditLog()

    @async def execute(self, agent, task):
        """Execute agent in secure sandbox"""

        # Log the task
        self.audit_log.log_task_start(agent.id, task)

        try:
            # Run in sandbox
            result = await self.sandbox.run(agent, task)

            # Validate output
            self.validate_output(result)

            # Log completion
            self.audit_log.log_task_complete(agent.id, result)

        return result

    except Exception as e:
        self.audit_log.log_task_error(agent.id, e)
        raise

```

## 8.9 Key Takeaways

1. Coding agents have evolved from autocomplete to autonomous development assistants.

2. **Multi-agent architectures** mirror development teams, with specialized agents for architecture, implementation, testing, and review.
  3. **AGENTS.md** is the emerging standard for providing agents with project-specific instructions.
  4. **Scaffolding** for coding agents includes context management, tool registries, and security boundaries.
  5. **Human review** remains essential—agents create PRs for review, not direct commits.
  6. **Incremental changes** with continuous validation are safer than large rewrites.
  7. **Security** must be designed in from the start: sandboxing, permissions, and audit logging.
-

# Chapter 9

# Agents for Mathematics and Physics

## 9.1 Chapter Preview

- Identify where formal methods and CAS tools fit in agent workflows.
- Separate runnable tooling from conceptual pseudocode.
- Highlight verification pitfalls and how to avoid them.

**Note:** Many examples in this chapter are illustrative pseudocode unless explicitly labeled as runnable, because formal tooling and CAS systems require environment-specific setup.

## 9.2 Introduction

Mathematics and physics present unique challenges for AI agents. Unlike coding, where correctness can often be verified through tests, mathematical reasoning requires formal proof and physical models demand empirical validation. This chapter explores specialized agents for scientific domains, their architectures, and the scaffolding required to support rigorous reasoning.

## 9.3 The Landscape of Scientific Agents

### 9.3.1 Distinct Requirements

Scientific agents differ from coding agents in several key ways:

| Aspect              | Coding Agents          | Scientific Agents                      |
|---------------------|------------------------|--|
| <b>Verification</b> | Tests, linting         | Formal proofs, experimental validation |
| <b>Precision</b>    | Functional correctness | Mathematical rigor                     |
| <b>Output</b>       | Source code            | Theorems, proofs, equations            |
| <b>Tools</b>        | IDEs, compilers        | Proof assistants, CAS, simulators      |
| <b>Context</b>      | Codebase               | Theorems, papers, datasets             |

### 9.3.2 Categories of Scientific Agents

1. **Theorem Proving Agents:** Construct formal proofs in systems like Lean (<https://lean-lang.org/learn/>), Coq (<https://coq.inria.fr/>), or Isabelle (<https://isabelle.in.tum.de/>)
2. **Symbolic Computation Agents:** Work with computer algebra systems (CAS)
3. **Numerical Simulation Agents:** Set up and run physics simulations
4. **Research Assistants:** Search literature, summarize findings, identify gaps

5. Educational Scaffolding Agents: Help students learn mathematical and physical concepts

## 9.4 Theorem Proving Agents

### 9.4.1 Formal Verification Background

Formal theorem proving ensures mathematical correctness through rigorous logical derivation. Unlike informal proofs in papers, formal proofs are machine-verifiable.

### 9.4.2 Ax-Prover Architecture

**Note:** Ax-Prover is a hypothetical composite example used to illustrate multi-agent theorem proving patterns.

```
class AxProverAgent:
    """Multi-agent theorem proving architecture inspired by Ax-Prover"""

    def __init__(self, llm, proof_assistant):
        self.llm = llm
        self.proof_assistant = proof_assistant # e.g., Lean, Coq
        self.strategy_agents = {
            'decomposition': DecompositionAgent(llm),
            'lemma_search': LemmaSearchAgent(llm),
            'tactic_selection': TacticSelectionAgent(llm),
            'creativity': CreativityAgent(llm)
        }

    async def prove(self, theorem: str) -> ProofResult:
        """Attempt to prove a theorem"""

        # 1. Formalize the statement
        formal_statement = await self.formalize(theorem)

        # 2. Decompose into subgoals
        subgoals = await self.strategy_agents['decomposition'].decompose(
            formal_statement
        )

        # 3. Search for relevant lemmas
        lemmas = await self.strategy_agents['lemma_search'].search(
            formal_statement, subgoals
        )

        # 4. Generate proof attempts
        proof_attempts = await self.generate_proof_attempts(
            formal_statement, subgoals, lemmas
        )

        # 5. Verify with proof assistant
        for attempt in proof_attempts:
            result = await self.proof_assistant.check(attempt)
            if result.verified:
                return ProofResult(success=True, proof=attempt)

        return ProofResult(success=False, partial_proofs=proof_attempts)

    async def formalize(self, natural_language: str) -> str:
        """Convert natural language to formal notation"""

```

```

prompt = f"""
Convert the following mathematical statement to formal Lean 4 syntax:

Statement: {natural_language}

Provide the formal statement only.
"""

return await self.llm.generate(prompt)

```

### 9.4.3 Integration with Proof Assistants

Agents connect to proof assistants through well-defined interfaces:

```

class LeanProofAssistant:
    """Interface to Lean 4 proof assistant"""

    def __init__(self, project_path: str):
        self.project_path = project_path
        self.server = LeanServer(project_path)

    async def check(self, proof: str) -> VerificationResult:
        """Verify a proof in Lean"""

        # Write proof to file
        proof_file = self.write_proof(proof)

        # Run Lean verification
        result = await self.server.check_file(proof_file)

        return VerificationResult(
            verified=not result.has_errors,
            errors=result.errors,
            goals=result.remaining_goals
        )

    async def get_available_tactics(self, goal_state: str) -> list:
        """Get tactics applicable to current goal state"""
        return await self.server.suggest_tactics(goal_state)

    async def search_mathlib(self, query: str) -> list:
        """Search Mathlib for relevant lemmas"""
        return await self.server.library_search(query)

```

### 9.4.4 Challenges in Theorem Proving

1. **Search Space Explosion:** Proofs can have many possible paths
2. **Creativity Required:** Non-obvious proof strategies
3. **Formalization Gap:** Translating informal to formal
4. **Domain Knowledge:** Deep mathematical understanding needed

## 9.5 Symbolic Computation Agents

### 9.5.1 Computer Algebra Systems

Symbolic computation agents work with systems like Mathematica, SymPy, or SageMath:

```

class SymbolicComputationAgent:
    """Agent for symbolic mathematical computation"""

    def __init__(self, llm, cas_backend='sympy'):
        self.llm = llm
        self.cas = self.initialize_cas(cas_backend)

    async def solve(self, problem: str) -> Solution:
        """Solve a mathematical problem symbolically"""

        # 1. Parse the problem
        parsed = await self.parse_problem(problem)

        # 2. Identify the type of problem
        problem_type = await self.classify_problem(parsed)

        # 3. Select appropriate methods
        methods = self.get_methods(problem_type)

        # 4. Attempt solutions
        for method in methods:
            try:
                result = await self.apply_method(method, parsed)
                if result.is_valid:
                    return Solution(
                        answer=result.answer,
                        method=method,
                        steps=result.steps
                    )
            except ComputationError:
                continue

        return Solution(success=False, attempted_methods=methods)

    async def simplify(self, expression: str) -> str:
        """Simplify a mathematical expression"""

        # Convert to CAS format
        cas_expr = self.cas.parse(expression)

        # Apply simplification
        simplified = self.cas.simplify(cas_expr)

        # Convert back to readable format
        return self.cas.to_latex(simplified)

    async def compute_integral(self, integrand: str, variable: str,
                             bounds: tuple = None) -> str:
        """Compute definite or indefinite integral"""

        expr = self.cas.parse(integrand)
        var = self.cas.symbol(variable)

        if bounds:
            result = self.cas.integrate(expr, (var, bounds[0], bounds[1]))
        else:
            result = self.cas.integrate(expr, var)

```

```
    return self.cas.to_latex(result)
```

### 9.5.2 Combining Symbolic and Neural Approaches

Modern agents combine symbolic precision with neural flexibility:

```
class HybridMathAgent:  
    """Combine symbolic computation with LLM reasoning"""  
  
    def __init__(self, llm, cas):  
        self.llm = llm  
        self.cas = cas  
  
    async def solve_with_explanation(self, problem: str) -> dict:  
        """Solve and explain a mathematical problem"""  
  
        # 1. LLM plans the solution strategy  
        strategy = await self.llm.generate(f"""  
        Given this problem: {problem}  
  
        Outline a step-by-step solution strategy.  
        Identify which steps require symbolic computation.  
        """)  
  
        # 2. Parse strategy into executable steps  
        steps = self.parse_strategy(strategy)  
  
        # 3. Execute each step  
        results = []  
        for step in steps:  
            if step.requires_symbolic:  
                result = await self.cas_execute(step)  
            else:  
                result = await self.llm_execute(step)  
            results.append(result)  
  
        # 4. Compile final answer with explanation  
        return {  
            'answer': results[-1],  
            'steps': results,  
            'explanation': await self.generate_explanation(results)  
        }
```

## 9.6 Physics Simulation Agents

### 9.6.1 Computational Physics Workflows

Physics agents orchestrate simulation workflows:

```
class PhysicsSimulationAgent:  
    """Agent for physics simulations"""  
  
    def __init__(self, llm, simulators):  
        self.llm = llm  
        self.simulators = {  
            'molecular_dynamics': MDSimulator(),  
            'quantum': QMSimulator(),
```

```

    'classical': ClassicalSimulator(),
    'fluid': CFDSimulator()
}

async def run_simulation(self, description: str) -> SimulationResult:
    """Set up and run a physics simulation from natural language"""

    # 1. Understand the physical system
    system_spec = await self.understand_system(description)

    # 2. Select appropriate simulator
    simulator = self.select_simulator(system_spec)

    # 3. Generate simulation parameters
    params = await self.generate_parameters(system_spec)

    # 4. Validate physical consistency
    await self.validate_physics(params)

    # 5. Run simulation
    result = await simulator.run(params)

    # 6. Analyze results
    analysis = await self.analyze_results(result, system_spec)

    return SimulationResult(
        raw_data=result,
        analysis=analysis,
        visualizations=await self.generate_plots(result)
    )

async def validate_physics(self, params: dict):
    """Ensure simulation parameters are physically consistent"""

    # Check conservation laws
    if not self.check_energy_conservation(params):
        raise PhysicsError("Energy conservation violated")

    # Check dimensional consistency
    if not self.check_dimensions(params):
        raise PhysicsError("Dimensional inconsistency")

    # Check boundary conditions
    if not self.check_boundaries(params):
        raise PhysicsError("Invalid boundary conditions")

```

## 9.6.2 Quantum Physics Specialization

Quantum physics requires specialized handling:

```

class QuantumPhysicsAgent:
    """Specialized agent for quantum mechanical problems"""

    def __init__(self, llm, qm_tools):
        self.llm = llm
        self.tools = qm_tools

    async def solve_schrodinger(self, system: str) -> dict:

```

```

"""Solve Schrödinger equation for a system"""

# 1. Construct Hamiltonian
hamiltonian = await self.construct_hamiltonian(system)

# 2. Identify symmetries
symmetries = await self.find_symmetries(hamiltonian)

# 3. Choose solution method
method = self.select_method(hamiltonian, symmetries)

# 4. Solve
if method == 'analytical':
    solution = await self.analytical_solve(hamiltonian)
elif method == 'numerical':
    solution = await self.numerical_solve(hamiltonian)
elif method == 'variational':
    solution = await self.variational_solve(hamiltonian)

return {
    'eigenstates': solution.states,
    'eigenvalues': solution.energies,
    'method': method,
    'symmetries': symmetries
}

async def compute_observable(self, state, observable: str) -> complex:
    """Compute expectation value of an observable"""

    operator = await self.construct_operator(observable)
    return await self.tools.expectation_value(state, operator)

```

## 9.7 Scaffolding for Scientific Agents

### 9.7.1 Tool Integration Layer

Scientific agents need access to specialized tools:

```

# Scientific agent tool configuration
tools:
    proof_assistants:
        lean4:
            path: /usr/local/bin/lean
            mathlib_path: ~/.elan/toolchains/leanprover--lean4---v4.3.0/lib/lean4/library
        coq:
            path: /usr/bin/coqc

    computer_algebra:
        sympy:
            module: sympy
        mathematica:
            path: /usr/local/bin/WolframScript

    simulators:
        molecular_dynamics:
            backend: lammps
            path: /usr/bin/lmp
        quantum:

```

```

backend: qiskit

visualization:
    matplotlib: true
    plotly: true
    manim: true

```

### 9.7.2 Knowledge Base Integration

Scientific agents need access to mathematical knowledge:

```

class MathematicalKnowledgeBase:
    """Knowledge base for mathematical agents"""

    def __init__(self):
        self.theorem_database = TheoremDatabase()
        self.formula_index = FormulaIndex()
        self.paper_embeddings = PaperEmbeddings()

    async def search_theorems(self, query: str) -> list:
        """Search for relevant theorems"""

        # Semantic search over theorem statements
        results = await self.theorem_database.semantic_search(query)

        # Include related lemmas and corollaries
        expanded = []
        for theorem in results:
            expanded.append(theorem)
            expanded.extend(await self.get_related(theorem))

        return expanded

    async def get_formula(self, name: str) -> Formula:
        """Retrieve a named formula"""
        return await self.formula_index.get(name)

    async def search_literature(self, topic: str) -> list:
        """Search mathematical literature"""

        # Search arXiv, Mathlib docs, textbooks
        papers = await self.paper_embeddings.search(topic)
        return papers

```

### 9.7.3 Verification Pipeline

All scientific agent outputs should be verified:

```

class ScientificVerificationPipeline:
    """Verify correctness of scientific agent outputs"""

    def __init__(self):
        self.proof_checker = ProofChecker()
        selfdimensional_analyzer = DimensionalAnalyzer()
        selfnumerical_validator = NumericalValidator()

    async def verify(self, output: ScientificOutput) -> VerificationResult:
        """Verify scientific output for correctness"""

```

```

checks = []

# 1. Check formal proofs
if output.has_proofs:
    proof_check = await self.proof_checker.verify(output.proofs)
    checks.append(('proofs', proof_check))

# 2. Check dimensional consistency
if output.has_equations:
    dim_check = await selfdimensional_analyzer.check(output.equations)
    checks.append(('dimensions', dim_check))

# 3. Numerical validation
if output.has_computations:
    num_check = await selfnumerical_validator.validate(
        output.computations
    )
    checks.append(('numerical', num_check))

# 4. Cross-check with known results
known_check = await self.check_against_known(output)
checks.append(('known_results', known_check))

return VerificationResult(
    verified=all(c[1].passed for c in checks),
    checks=checks
)

```

## 9.8 Educational Scaffolding Agents

### 9.8.1 Mathematics Education

AI agents are transforming mathematics education:

```

class MathTutoringAgent:
    """Agent for mathematics education and tutoring"""

    def __init__(self, llm, level='undergraduate'):
        self.llm = llm
        self.level = level
        self.student_model = StudentModel()

    @async def explain_concept(self, concept: str) -> str:
        """Explain a mathematical concept at appropriate level"""

        # Get student's current understanding
        background = await self.student_model.get_background()

        # Generate explanation
        explanation = await self.llm.generate(f"""
            Explain {concept} to a student with this background: {background}

            Level: {self.level}

            Include:
            - Intuitive explanation
            - Formal definition
        """)

```

```

        - Key examples
        - Common misconceptions
        - Connection to prior knowledge
        """)

    return explanation

async def generate_problems(self, topic: str, count: int,
                           difficulty: str) -> list:
    """Generate practice problems with solutions"""

    problems = await self.llm.generate(f"""
Generate {count} {difficulty} problems on {topic}.

For each problem provide:
1. Problem statement
2. Hints (progressive)
3. Complete solution
4. Common errors to avoid
""")

    return self.parse_problems(problems)

async def provide_feedback(self, student_work: str,
                           problem: str) -> Feedback:
    """Analyze student work and provide feedback"""

    analysis = await self.llm.generate(f"""
Analyze this student's solution:

Problem: {problem}
Student work: {student_work}

Provide:
1. Is the final answer correct?
2. Are the intermediate steps correct?
3. What misconceptions are evident?
4. Specific suggestions for improvement
5. Encouragement and next steps
""")

    return self.parse_feedback(analysis)

```

## 9.8.2 Physics Education

Physics scaffolding addresses visualization challenges:

```

class PhysicsEducationAgent:
    """Agent for physics education with visualization"""

    def __init__(self, llm, visualizer):
        self.llm = llm
        self.visualizer = visualizer

    async def explain_with_simulation(self, concept: str) -> dict:
        """Explain physics concept with interactive simulation"""

        # Generate explanation

```

```

explanation = await self.explain_concept(concept)

# Create visualization parameters
viz_params = await self.generate_visualization_params(concept)

# Generate simulation
simulation = await self.visualizer.create_simulation(viz_params)

# Create interactive exploration tasks
tasks = await self.generate_exploration_tasks(concept)

return {
    'explanation': explanation,
    'simulation': simulation,
    'exploration_tasks': tasks,
    'key_parameters': viz_params['adjustable']
}

async def analyze_misconception(self, student_statement: str) -> dict:
    """Identify and address physics misconceptions"""

    analysis = await self.llm.generate(f"""
The student said: "{student_statement}"

1. Identify any physics misconceptions
2. Explain the correct physics
3. Suggest experiments or simulations to demonstrate
4. Provide an analogy that builds correct intuition
""")

    return self.parse_misconception_analysis(analysis)

```

## 9.9 Research Agent Workflows

### 9.9.1 Literature Review Agents

Agents that assist with scientific literature:

```

class LiteratureReviewAgent:
    """Agent for mathematical and physics literature review"""

    def __init__(self, llm, databases):
        self.llm = llm
        self.databases = {
            'arxiv': ArxivAPI(),
            'mathscinet': MathSciNetAPI(),
            'semantic_scholar': SemanticScholarAPI()
        }

    async def survey_topic(self, topic: str) -> Survey:
        """Create a survey of a research topic"""

        # 1. Search for relevant papers
        papers = await self.search_all_databases(topic)

        # 2. Cluster by approach/contribution
        clusters = await self.cluster_papers(papers)

```

```

# 3. Identify key results
key_results = await self.extract_key_results(papers)

# 4. Find open problems
open_problems = await self.identify_open_problems(papers)

# 5. Generate survey
survey = await self.generate_survey(
    clusters, key_results, open_problems
)

return survey

async def find_related_work(self, paper_or_idea: str) -> list:
    """Find work related to a paper or research idea"""

    # Extract key concepts
    concepts = await self.extract_concepts(paper_or_idea)

    # Search for related papers
    related = []
    for concept in concepts:
        papers = await self.search_concept(concept)
        related.extend(papers)

    # Rank by relevance
    ranked = await self.rank_relevance(related, paper_or_idea)

    return ranked[:20] # Top 20 most relevant

```

## 9.10 Best Practices

### 9.10.1 Rigorous Verification

Always verify scientific outputs:

```

async def execute_with_verification(self, task):
    result = await self.agent.execute(task)

    # Verify before returning
    verification = await self.verifier.verify(result)

    if not verification.passed:
        raise VerificationError(
            f"Output failed verification: {verification.errors}"
        )

    return result

```

### 9.10.2 Uncertainty Quantification

Scientific agents should express uncertainty:

```

class UncertaintyAwareAgent:
    """Agent that quantifies uncertainty in results"""

    async def solve(self, problem):

```

```

        result = await self.compute(problem)

        # Quantify uncertainty
        uncertainty = await self.estimate_uncertainty(result, problem)

    return {
        'result': result,
        'uncertainty': uncertainty,
        'confidence': self.compute_confidence(uncertainty)
    }

```

### 9.10.3 Reproducibility

Ensure all computations are reproducible:

```

class ReproducibleComputation:
    """Ensure scientific computations are reproducible"""

    def __init__(self):
        self.rng_seed = None
        self.version_info = {}

    def setup(self, seed: int):
        """Set up reproducible environment"""
        self.rng_seed = seed
        np.random.seed(seed)
        random.seed(seed)

        # Record versions
        self.version_info = {
            'numpy': np.__version__,
            'scipy': scipy.__version__,
            'python': sys.version
        }

    def get_reproduction_info(self):
        """Get information needed to reproduce computation"""
        return {
            'seed': self.rng_seed,
            'versions': self.version_info,
            'timestamp': datetime.now().isoformat()
        }

```

### 9.10.4 Domain Expert Collaboration

Design agents to work with domain experts:

```

class CollaborativeAgent:
    """Agent designed for collaboration with human experts"""

    async def propose_approach(self, problem):
        """Propose approach for expert review"""

        approaches = await self.generate_approaches(problem)

    return {
        'approaches': approaches,
        'recommendation': approaches[0],

```

```

        'rationale': await self.explain_recommendation(approaches[0]),
        'request_for_feedback': True
    }

async def incorporate_feedback(self, feedback, current_state):
    """Incorporate expert feedback into solution process"""

    # Parse feedback
    parsed = await self.parse_expert_feedback(feedback)

    # Adjust approach
    adjusted = await self.adjust_approach(current_state, parsed)

    return adjusted

```

## 9.11 Key Takeaways

1. **Scientific agents** require formal verification and rigorous validation beyond what coding agents need.
  2. **Theorem proving agents** combine LLM creativity with proof assistant verification for mathematical rigor.
  3. **Symbolic computation** and neural approaches are complementary—use both for best results.
  4. **Physics agents** must respect conservation laws, dimensional consistency, and physical constraints.
  5. **Educational scaffolding** agents adapt explanations to student level and address misconceptions.
  6. **Verification pipelines** should check proofs, dimensions, and compare with known results.
  7. **Reproducibility** is essential—record seeds, versions, and all parameters.
  8. **Collaboration** with domain experts improves agent reliability and trustworthiness.
-

# Chapter 10

# Common Failure Modes, Testing, and Fixes

## 10.1 Chapter Goals

By the end of this chapter, you should be able to:

- Recognize the most common ways agentic workflows fail in production.
- Design a test strategy that catches failures before deployment.
- Apply practical mitigation and recovery patterns.
- Turn incidents into durable process and architecture improvements.

## 10.2 Why Failures Are Different in Agentic Systems

Traditional software failures are often deterministic and reproducible. Agent failures can also include:

- **Nondeterminism** from model sampling and external context.
- **Tool and API variance** across environments and versions.
- **Instruction ambiguity** when prompts, policy files, or skills conflict.
- **Long-horizon drift** where behavior degrades over many steps.

This means reliability work must combine classic software testing with scenario-based evaluation and operational controls.

## 10.3 Failure Taxonomy

Use this taxonomy to classify incidents quickly and choose the right fix path.

### 10.3.1 1) Planning and Reasoning Failures

Symptoms:

- Agent picks the wrong sub-goal.
- Repeats work or loops without convergence.
- Produces plausible but invalid conclusions.

Typical causes:

- Missing constraints in system instructions.
- Overly broad tasks with no decomposition guardrails.

- No termination criteria.

Fast fixes:

- Add explicit success criteria and stop conditions.
- Break tasks into bounded steps.
- Require intermediate checks before irreversible actions.

### **10.3.2 2) Tooling and Integration Failures**

Symptoms:

- Tool calls fail intermittently.
- Wrong parameters passed to tools.
- Tool output parsed incorrectly.

Typical causes:

- Schema drift or undocumented API changes.
- Weak input validation.
- Inconsistent retry/backoff handling.

Fast fixes:

- Validate tool contracts at runtime.
- Add strict argument schemas.
- Standardize retries with idempotency keys.

### **10.3.3 3) Context and Memory Failures**

Symptoms:

- Agent forgets prior constraints.
- Important instructions are dropped when context grows.
- Stale memories override fresh data.

Typical causes:

- Context window pressure.
- Poor memory ranking/retrieval.
- Missing recency and source-quality weighting.

Fast fixes:

- Introduce context budgets and summarization checkpoints.
- Add citation requirements for retrieved facts.
- Expire or down-rank stale memory entries.

### **10.3.4 4) Safety and Policy Failures**

Symptoms:

- Sensitive files modified unexpectedly.
- Security constraints bypassed through tool chains.
- Unsafe suggestions in generated code.

Typical causes:

- Weak policy enforcement boundaries.
- No pre-merge policy gates.
- Implicit trust in generated output.

Fast fixes:

- Enforce allow/deny lists at tool gateway level.
- Require policy checks in CI.
- Route high-risk actions through human approval.

### 10.3.5 5) Collaboration and Workflow Failures

Symptoms:

- Multiple agents make conflicting changes.
- PRs churn with contradictory edits.
- Work stalls due to unclear ownership.

Typical causes:

- Missing orchestration contracts.
- No lock/lease model for shared resources.
- Role overlap without clear handoff rules.

Fast fixes:

- Add ownership rules per path/component.
- Use optimistic locking with conflict resolution policy.
- Define role-specific done criteria.

## 10.4 Testing Strategy for Agentic Workflows

A robust strategy uses multiple test layers. No single test type is sufficient.

### 10.5 1. Static and Structural Checks

Use these to fail fast before expensive model execution:

- Markdown/schema validation for instruction files.
- Prompt template linting.
- Tool interface compatibility checks.
- Dependency and version constraint checks.

### 10.6 2. Deterministic Unit Tests (Without LLM Calls)

Test orchestration logic, parsers, and guards deterministically:

- State transitions.
- Retry and timeout behavior.
- Permission checks.
- Conflict resolution rules.

**Snippet status:** Runnable shape (simplified for clarity).

```
from dataclasses import dataclass

@dataclass
class StepResult:
    ok: bool
    retryable: bool
```

```

def should_retry(result: StepResult, attempt: int, max_attempts: int = 3) -> bool:
    return (not result.ok) and result.retryable and attempt < max_attempts

def test_retry_policy():
    assert should_retry(StepResult(ok=False, retryable=True), attempt=1)
    assert not should_retry(StepResult(ok=False, retryable=False), attempt=1)
    assert not should_retry(StepResult(ok=False, retryable=True), attempt=3)

```

## 10.7 3. Recorded Integration Tests (Golden Traces)

Capture representative interactions and replay them against newer builds:

- Record tool inputs/outputs.
- Freeze external dependencies where possible.
- Compare final artifacts and decision traces.

Use these to detect drift in behavior after prompt/tool/model changes.

## 10.8 4. Scenario and Adversarial Evaluations

Design “challenge suites” for known weak spots:

- Ambiguous requirements.
- Contradictory documentation.
- Missing dependencies.
- Partial outages and degraded APIs.
- Prompt-injection attempts in retrieved content.

Pass criteria should include not just correctness, but also:

- Policy compliance.
- Cost/latency ceilings.
- Evidence quality (citations, rationale quality).

## 10.9 5. Production Guardrail Tests

Before enabling autonomous writes/merges in production, validate:

- Protected-path enforcement.
- Secret scanning and license checks.
- Human approval routing for high-impact actions.
- Rollback path on failed deployments.

## 10.10 Practical Fix Patterns

When incidents happen, reusable fix patterns reduce MTTR (mean time to recovery).

### 10.10.1 Pattern A: Contract Hardening

- Add strict schemas between planner and tool runner.
- Reject malformed or out-of-policy requests early.
- Version contracts (v1, v2) and support migrations.

#### 10.10.2 Pattern B: Progressive Autonomy

- Start in “suggest-only” mode.
- Move to “execute with review” mode.
- Graduate to autonomous mode only after SLO compliance.

#### 10.10.3 Pattern C: Two-Phase Execution

1. **Plan phase:** generate proposed actions and expected effects.
2. **Apply phase:** execute only after policy and validation checks pass.

This reduces irreversible mistakes and improves auditability.

#### 10.10.4 Pattern D: Fallback and Circuit Breakers

- If tool failure rate spikes, disable affected paths automatically.
- Fall back to a safer baseline workflow.
- Alert operators with incident context.

#### 10.10.5 Pattern E: Human-in-the-Loop Escalation

Define explicit escalation triggers:

- Repeated retries without progress.
- Any request touching protected paths.
- Low-confidence output in high-risk domains.

### 10.11 Incident Response Runbook (Template)

Use a lightweight runbook so teams respond consistently:

1. **Detect:** Alert from CI, runtime monitor, or user report.
2. **Classify:** Map to taxonomy category.
3. **Contain:** Stop autonomous actions if blast radius is unclear.
4. **Diagnose:** Reproduce with trace + config snapshot.
5. **Mitigate:** Apply short-term guardrails/fallback.
6. **Fix:** Implement structural correction.
7. **Verify:** Re-run affected test suites + adversarial cases.
8. **Learn:** Add regression test and update docs.

### 10.12 Metrics That Actually Matter

Track these to evaluate reliability improvements over time:

- **Task success rate** (with policy compliance as part of success).
- **Intervention rate** (how often humans must correct the agent).
- **Escaped defect rate** (failures discovered after merge/deploy).
- **Mean time to detect (MTTD)** and **mean time to recover (MTTR)**.
- **Cost per successful task** and **latency percentiles**.

Avoid vanity metrics (for example, “number of agent runs”) without quality and safety context.

### 10.13 Anti-Patterns to Avoid

- Treating prompt edits as sufficient reliability work.
- Allowing autonomous writes without protected-path policies.

- Skipping regression suites after model/version upgrades.
- Relying on a single benchmark instead of diverse scenarios.
- Ignoring ambiguous ownership in multi-agent flows.

## 10.14 A Minimal Reliability Checklist

Before enabling broad production use, confirm:

- Snippets and examples are clearly labeled (Runnable/pseudocode/simplified).
- Tool contracts are versioned and validated.
- CI includes policy, security, and regression checks.
- Failure injection scenarios are part of routine testing.
- Rollback and escalation paths are documented and exercised.

## 10.15 Chapter Summary

Reliable agentic systems are built, not assumed. Teams that combine:

- clear contracts,
- layered testing,
- progressive autonomy,
- strong policy gates, and
- incident-driven learning

consistently outperform teams relying on prompt-only tuning.

In practice, your competitive advantage comes from how quickly you detect, contain, and permanently fix failures—not from avoiding them entirely.

# Bibliography

- GitHub Agentic Workflows documentation. <https://github.github.io/gh-aw/>. Accessed: 2026-02-05.
- GitHub Agentic Workflows repository. <https://github.com/github/gh-aw>. Accessed: 2026-02-05.
- GitHub Copilot documentation. <https://docs.github.com/en/copilot>. Accessed: 2026-02-05.
- GitHub Copilot coding agent. GitHub Docs. Accessed: 2026-02-05.
- Copilot coding agent environment customization. GitHub Docs. Accessed: 2026-02-05.
- Model Context Protocol (MCP). <https://modelcontextprotocol.io/>. Accessed: 2026-02-05.
- Skills Protocol documentation. <https://skillsprotocol.com/>. Accessed: 2026-02-05.
- Skills Protocol Implementation Guide. <https://skillsprotocol.com/implementation-guide>. Accessed: 2026-02-05.
- Skills Protocol specification. <https://skillsprotocol.com/specification>. Accessed: 2026-02-05.
- Agent Skills format: Skill structure. <https://skillsprotocol.com/skill-structure>. Accessed: 2026-02-05.
- Agent Skills format: Skill manifest. <https://skillsprotocol.com/skill-manifest>. Accessed: 2026-02-05.
- OpenAI. <https://openai.com/>. Accessed: 2026-02-05.
- Anthropic. <https://www.anthropic.com/>. Accessed: 2026-02-05.
- OpenAI Codex overview. <https://openai.com/index/introducing-codex/>. Accessed: 2026-02-05.
- OpenAI Codex skills documentation. <https://platform.openai.com/docs/guides/codex/skills>. Accessed: 2026-02-05.
- Claude Code documentation. <https://code.claude.com/docs>. Accessed: 2026-02-05.
- Cursor editor. <https://www.cursor.com/>. Accessed: 2026-02-05.
- CodeGPT. <https://codegpt.co/>. Accessed: 2026-02-05.
- LangChain documentation. <https://python.langchain.com/>. Accessed: 2026-02-05.
- LangGraph documentation. <https://langchain-ai.github.io/langgraph/>. Accessed: 2026-02-05.
- CrewAI documentation. <https://docs.crewai.com/>. Accessed: 2026-02-05.
- Microsoft Semantic Kernel. <https://learn.microsoft.com/semantic-kernel/>. Accessed: 2026-02-05.
- AutoGen documentation. <https://microsoft.github.io/autogen/>. Accessed: 2026-02-05.
- Ollama. <https://ollama.com/>. Accessed: 2026-02-05.
- Tailscale. <https://tailscale.com/>. Accessed: 2026-02-05.
- Lean documentation. <https://lean-lang.org/learn/>. Accessed: 2026-02-05.
- Coq. <https://coq.inria.fr/>. Accessed: 2026-02-05.
- Isabelle. <https://isabelle.in.tum.de/>. Accessed: 2026-02-05.